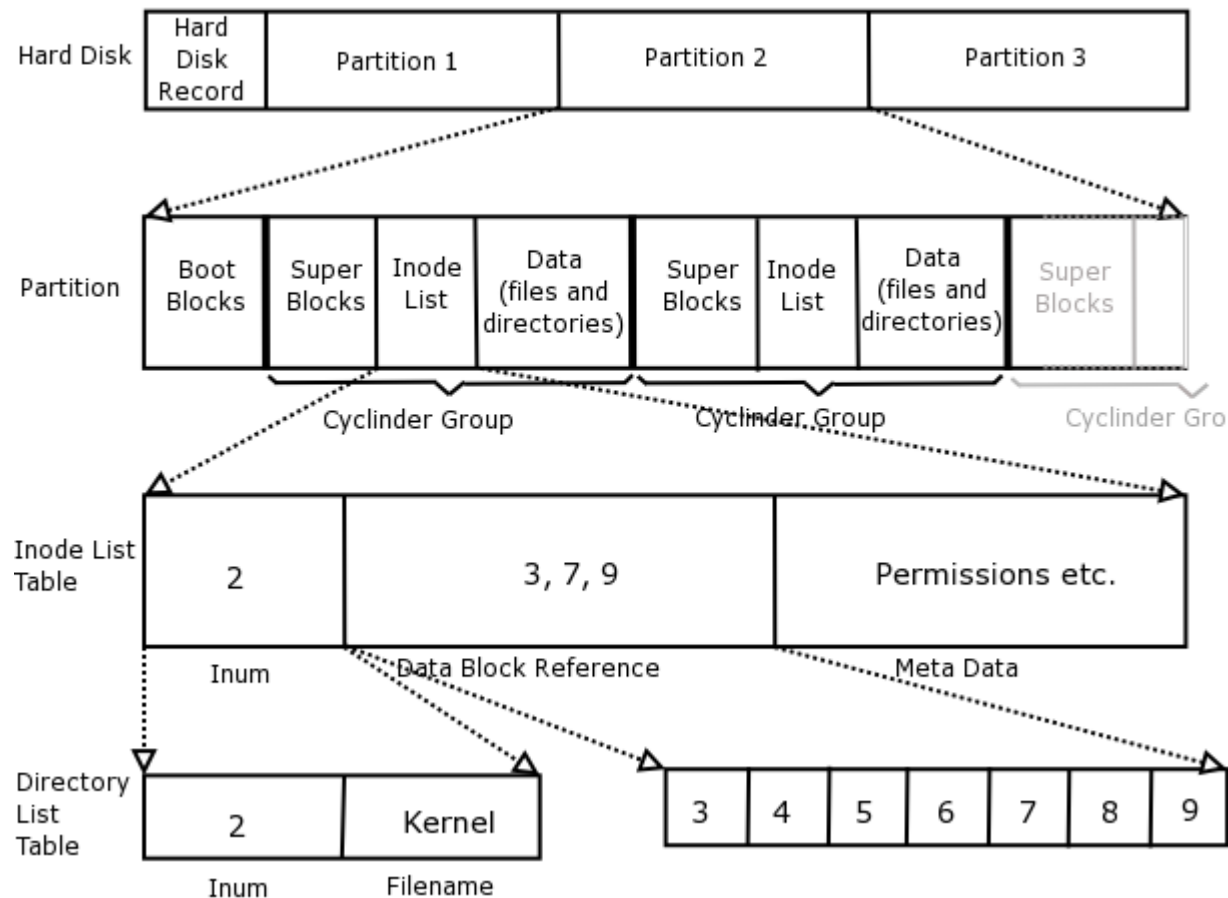


File Systems. Part 1

UNIX File System Layout



“... the abstractions of processes (and threads), address spaces, and **files are the most important concepts relating to operating systems.**

... **Files are logical units of information created by processes. A disk will usually contain thousands or even millions of them, each one independent of the others.**

... Processes can read existing files and create new ones if need be. Information stored in files must be **persistent, that is, not be affected by process creation and termination. A file should disappear only when its owner explicitly removes it.”**

“... Files are managed by the operating system.

How they are structured, named, accessed, used, protected, implemented, and managed are major topics in operating system design.

As a whole, that part of the operating system dealing with files is known as the **file system.”**

File Naming

“... A file is an abstraction mechanism. It provides a way to store information on the disk and read it back later. This must be done in such a way as to shield the user from the details of how and where the information is stored, and how the disks actually work.

Probably the most important characteristic of any abstraction mechanism is the way the objects being managed are named. ... When a process creates a file, it gives the file a name. When the process terminates, the file continues to exist and can be accessed by other processes using its name.”

Extension	Meaning
.bak	Backup file
.c	C source program
.gif	Compuserve Graphical Interchange Format image
.hlp	Help file
.html	World Wide Web HyperText Markup Language document
.jpg	Still picture encoded with the JPEG standard
.mp3	Music encoded in MPEG layer 3 audio format
.mpg	Movie encoded with the MPEG standard
.o	Object file (compiler output, not yet linked)
.pdf	Portable Document Format file
.ps	PostScript file
.tex	Input for the TEX formatting program
.txt	General text file
.zip	Compressed archive

Figure 4-1. Some typical file extensions.

File Structure

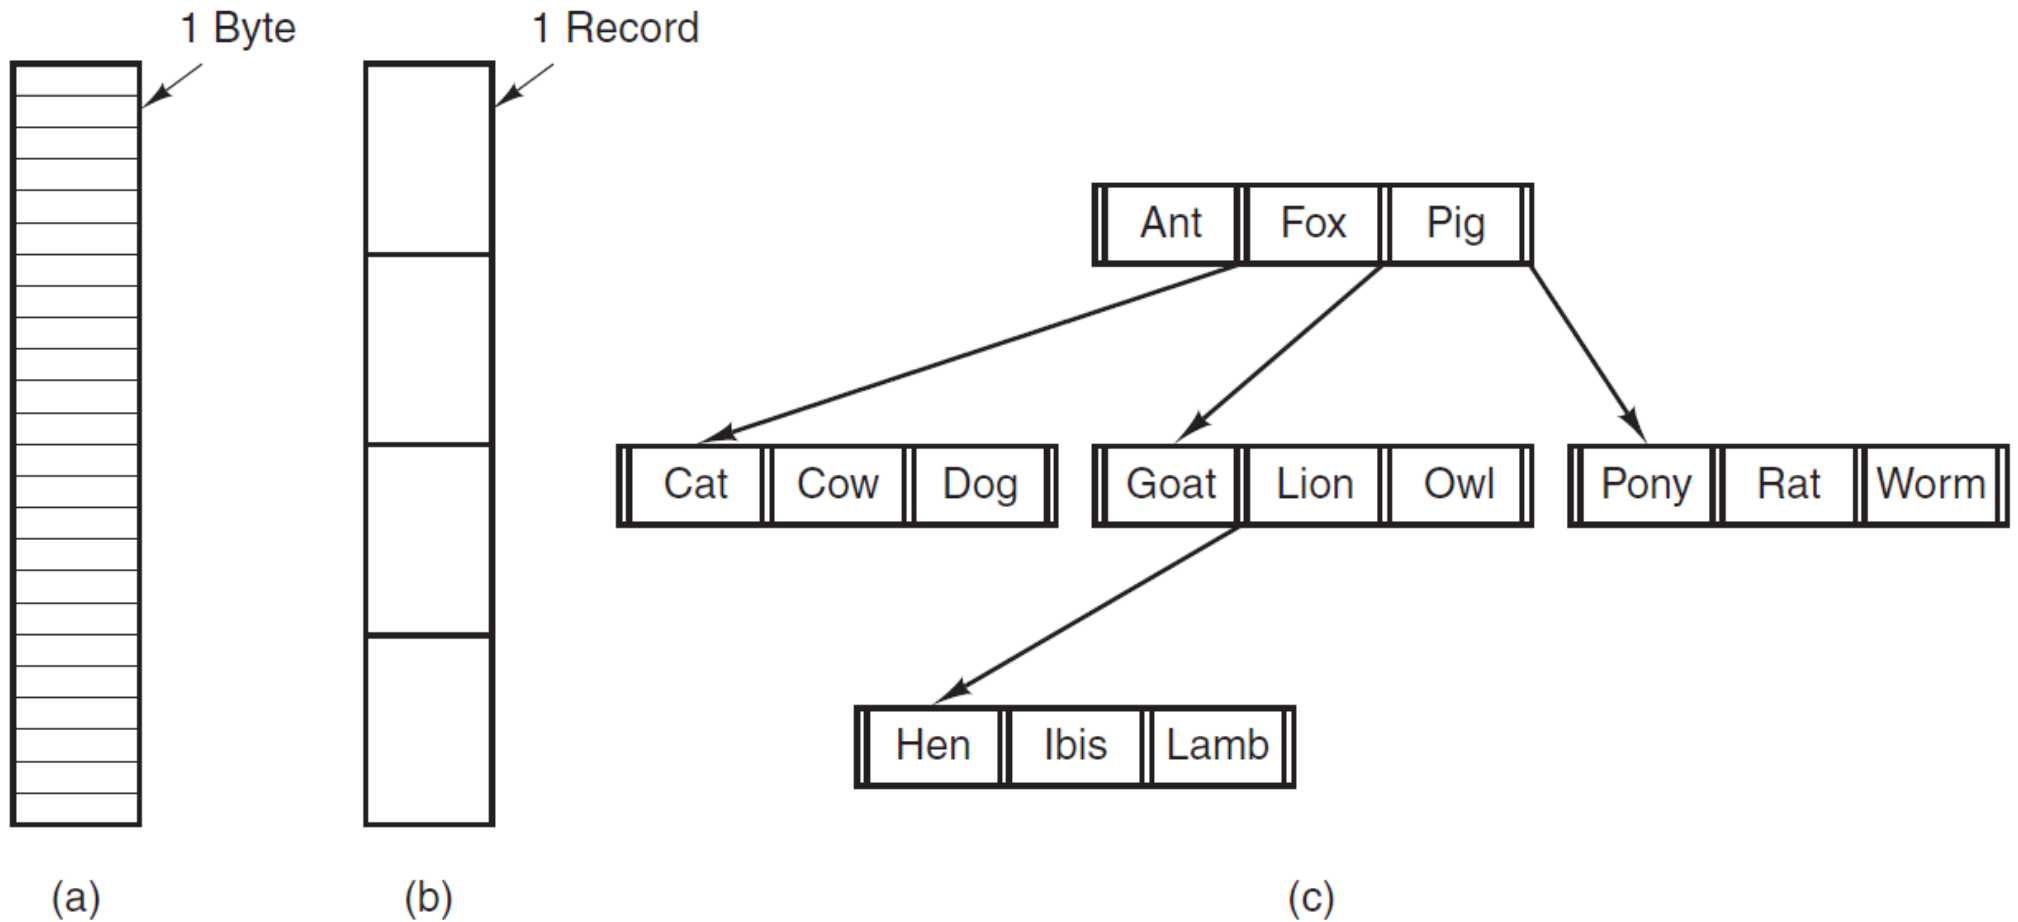


Figure 4-2. Three kinds of files. (a) Byte sequence. (b) Record sequence. (c) Tree.

File Types

Regular files

Directories

Character special files

Block special files

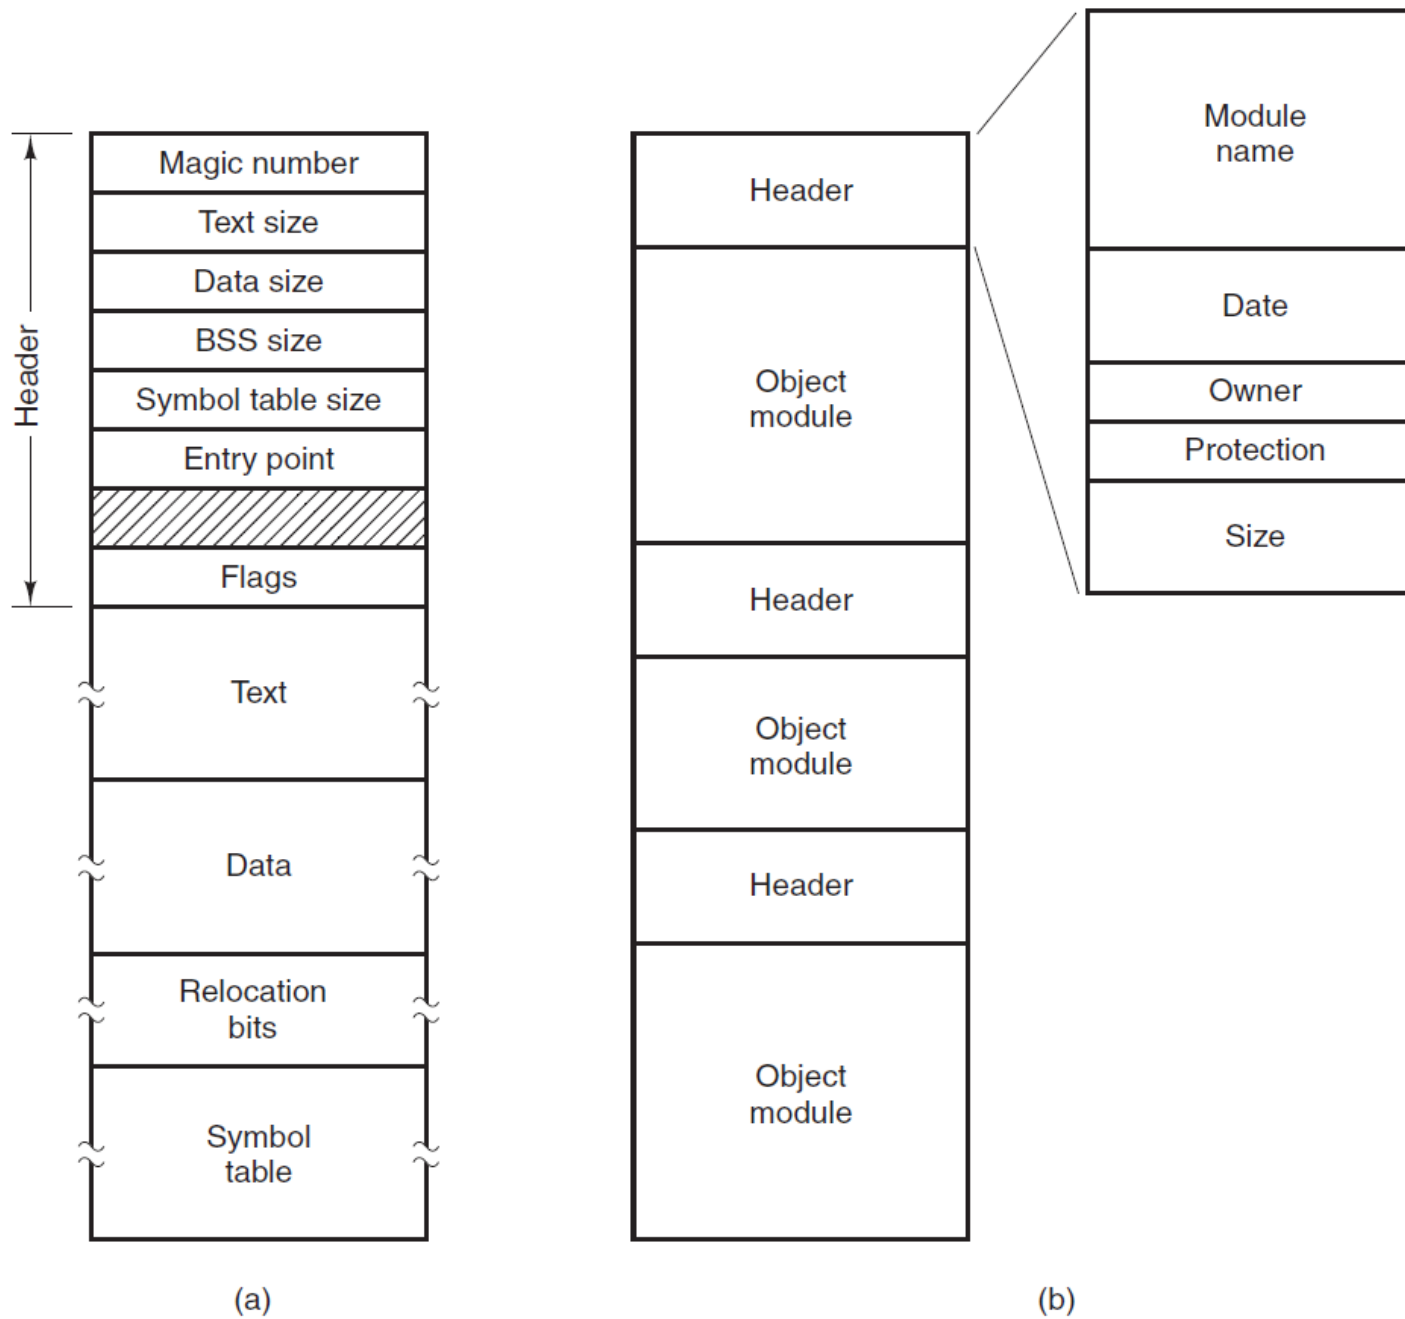


Figure 4-3. (a) An executable file. (b) An archive.

File Access

Sequential access

**Random access
operation seek**

Attribute	Meaning
Protection	Who can access the file and in what way
Password	Password needed to access the file
Creator	ID of the person who created the file
Owner	Current owner
Read-only flag	0 for read/write; 1 for read only
Hidden flag	0 for normal; 1 for do not display in listings
System flag	0 for normal files; 1 for system file
Archive flag	0 for has been backed up; 1 for needs to be backed up
ASCII/binary flag	0 for ASCII file; 1 for binary file
Random access flag	0 for sequential access only; 1 for random access
Temporary flag	0 for normal; 1 for delete file on process exit
Lock flags	0 for unlocked; nonzero for locked
Record length	Number of bytes in a record
Key position	Offset of the key within each record
Key length	Number of bytes in the key field
Creation time	Date and time the file was created
Time of last access	Date and time the file was last accessed
Time of last change	Date and time the file was last changed
Current size	Number of bytes in the file
Maximum size	Number of bytes the file may grow to

Figure 4-4. Some possible file attributes.

File Operations (system calls)

- 1. Create**
- 2. Delete**
- 3. Open**
- 4. Close**
- 5. Read**
- 6. Write**
- 7. Append**
- 8. Seek**
- 9. Get attributes**
- 10. Set attributes**
- 11. Rename**

/* File copy program. Error checking and reporting is minimal. */

```
#include <sys/types.h> /* include necessary header files */
#include <fcntl.h>
#include <stdlib.h>
#include <unistd.h>

int main(int argc, char *argv[]); /* ANSI prototype */

#define BUF_SIZE 4096 /* use a buffer size of 4096 bytes */
#define OUTPUT_MODE 0700 /* protection bits for output file */

int main(int argc, char *argv[])
{
    int in_fd, out_fd, rd_count, wt_count;
    char buffer[BUF_SIZE];

    if (argc != 3) exit(1); /* syntax error if argc is not 3 */

    /* Open the input file and create the output file */
    in_fd = open(argv[1], O_RDONLY); /* open the source file */
    if (in_fd < 0) exit(2); /* if it cannot be opened, exit */
    out_fd = creat(argv[2], OUTPUT_MODE); /* create the destination file */
    if (out_fd < 0) exit(3); /* if it cannot be created, exit */
```

```

/* Copy loop */
while (TRUE) {
    rd_count = read(in_fd, buffer, BUF_SIZE); /* read a block of data */
    if (rd_count <= 0) break;                  /* if end of file or error, exit loop */
    wt_count = write(out_fd, buffer, rd_count); /* write data */
    if (wt_count <= 0) exit(4);                /* wt_count <= 0 is an error */
}

/* Close the files */
close(in_fd);
close(out_fd);
if (rd_count == 0)                          /* no error on last read */
    exit(0);
else
    exit(5);                                /* error on last read */
}

```

Figure 4-5. A simple program to copy a file.

Directories:

Single-level Directory Systems




Figure 4-6. A single-level directory system containing four files.

MOS4E

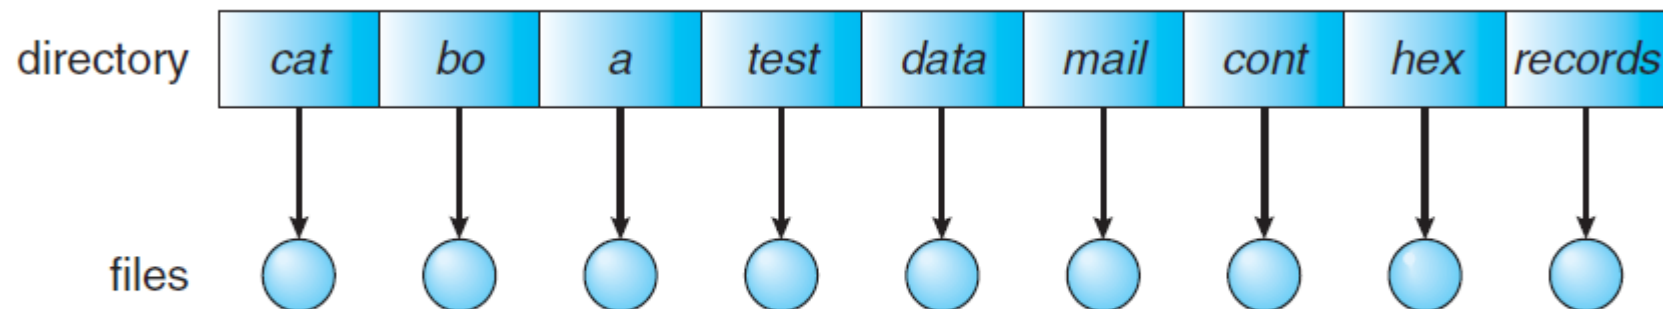


Figure 11.9 Single-level directory.

OSC9E

Directories:

Two-level Directory Systems

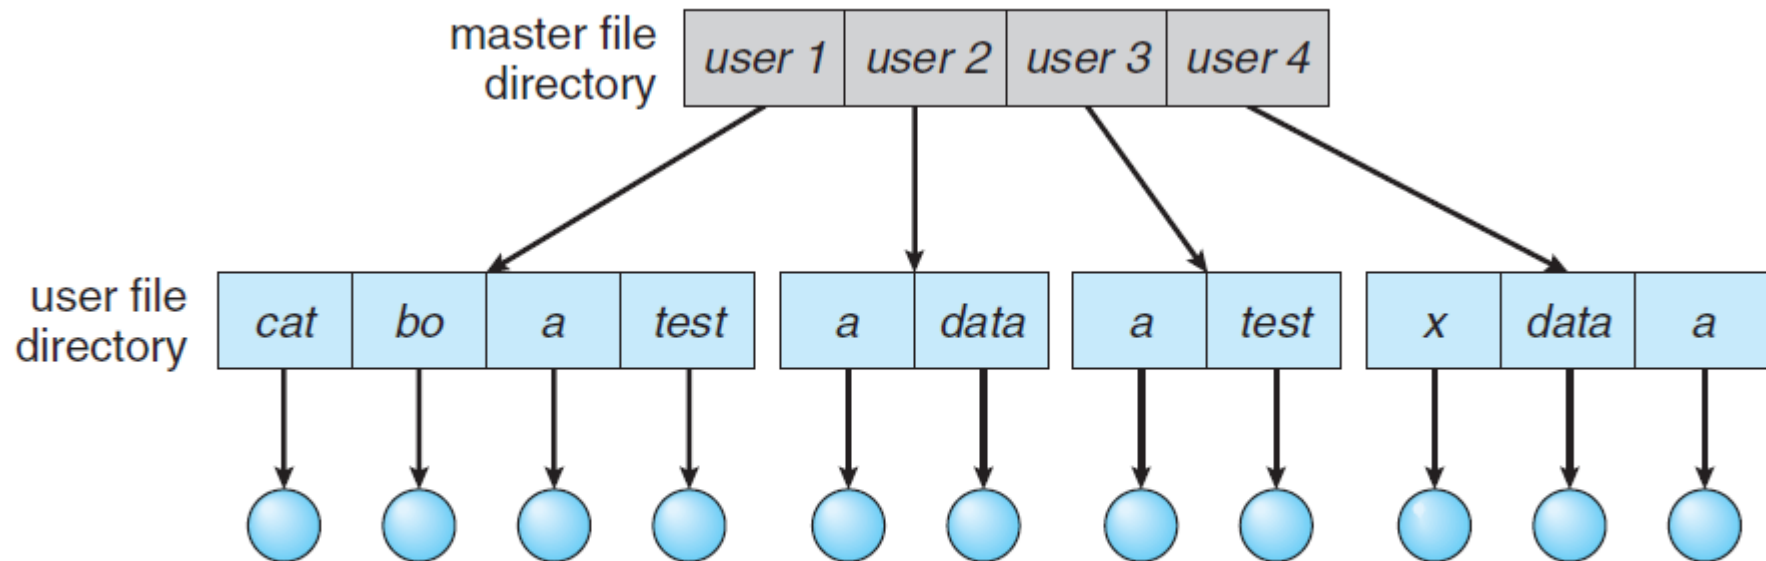


Figure 11.10 Two-level directory structure.

Directories:

Hierarchical Directory Systems

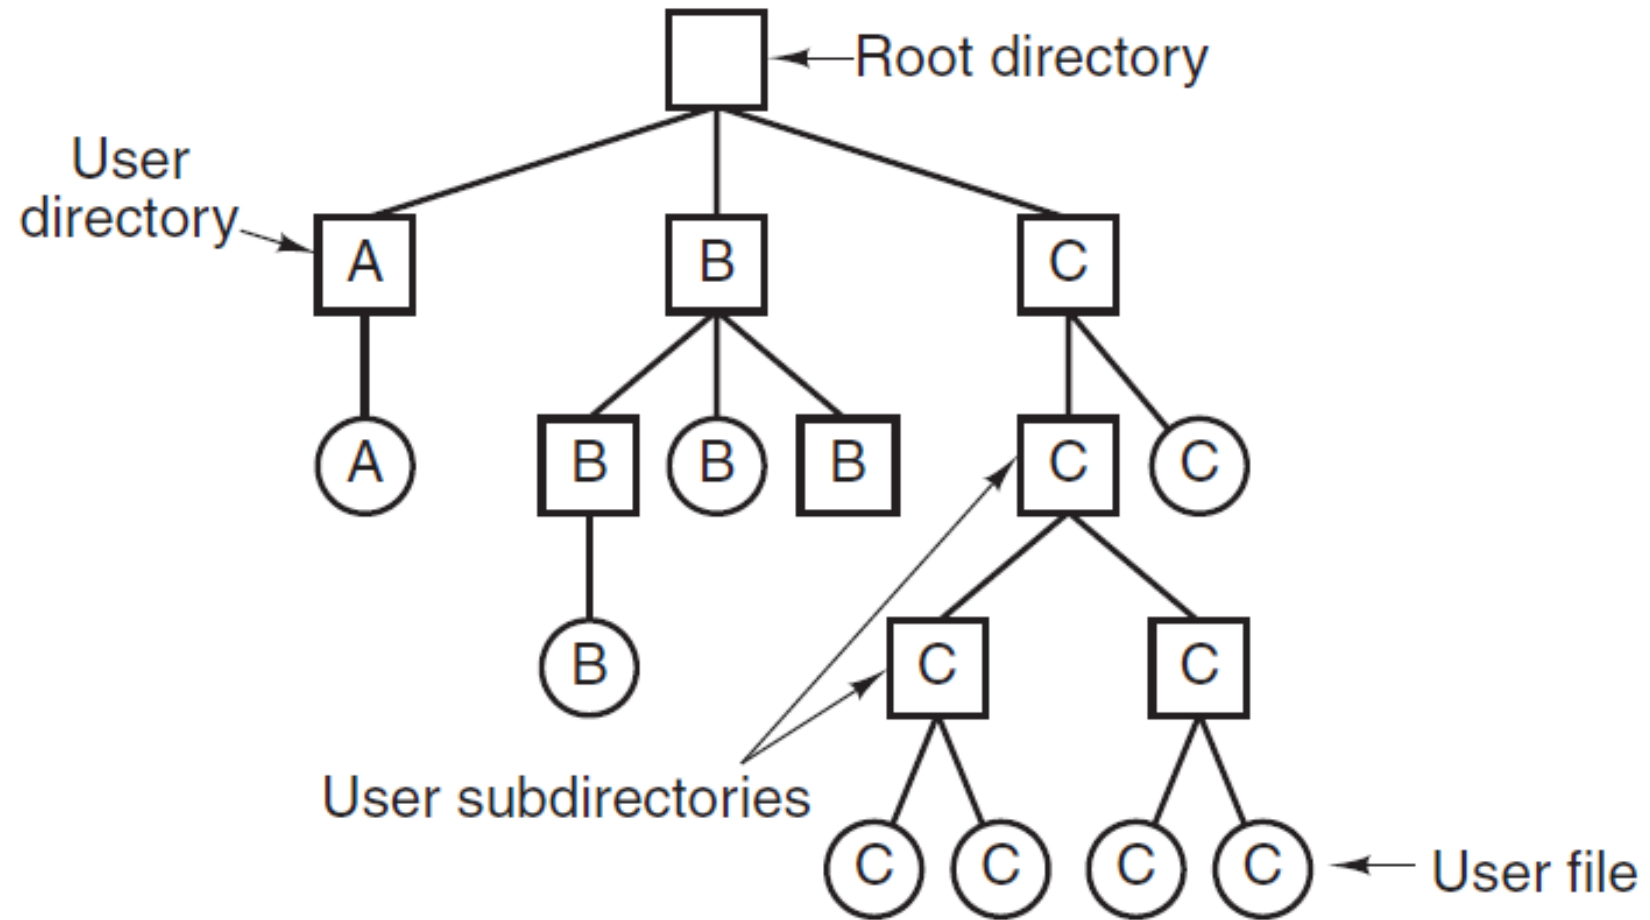


Figure 4-7. A hierarchical directory system.

Directories:

Hierarchical Directory Systems

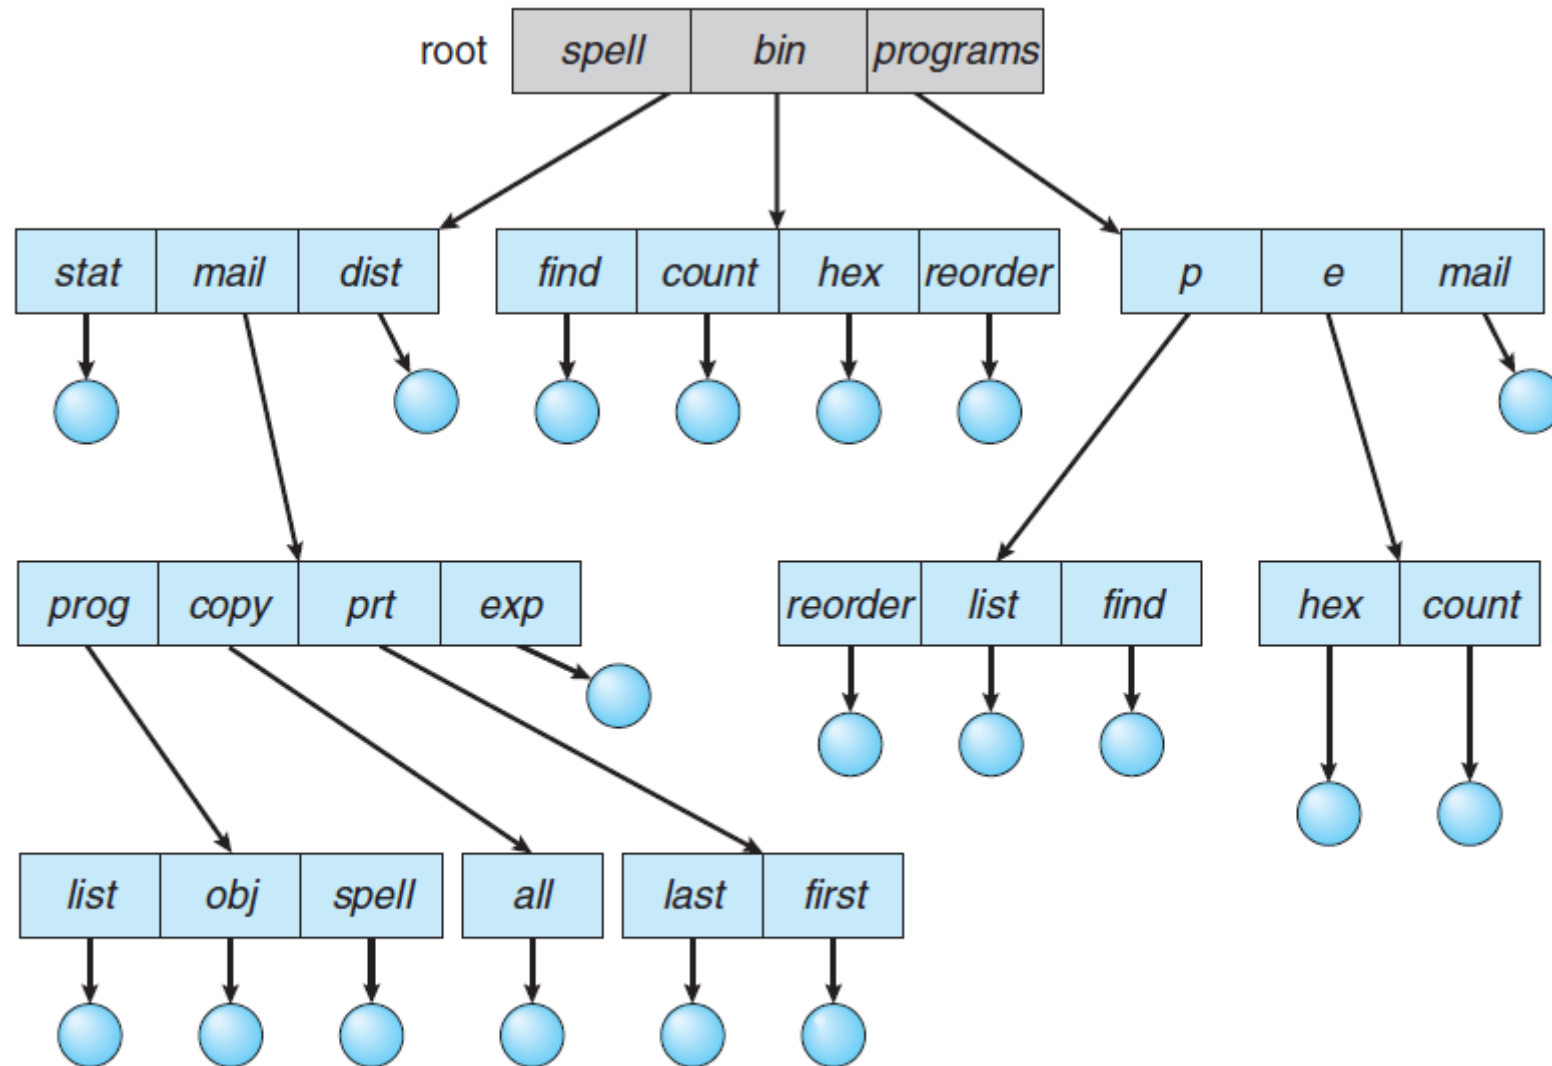


Figure 11.11 Tree-structured directory structure.

Path Names

Absolute path name

Relative path name

Working (current directory)

Special entries in every directory:

“.” (“dot”)

current directory

“..” (“dotdot”)

parent directory




Figure 4-8. A UNIX directory tree.

Directories:

Acyclic-Graph Directories

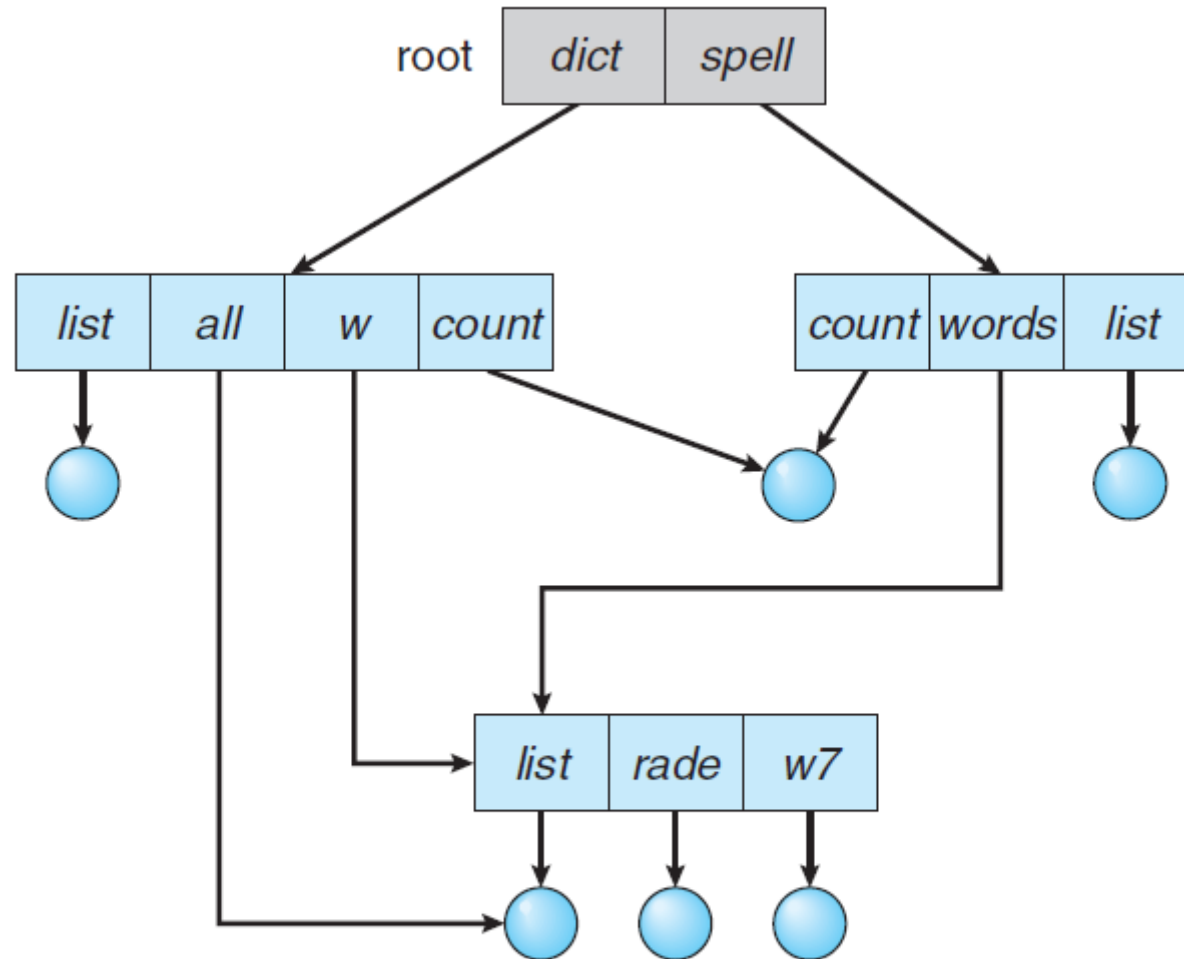


Figure 11.12 Acyclic-graph directory structure.

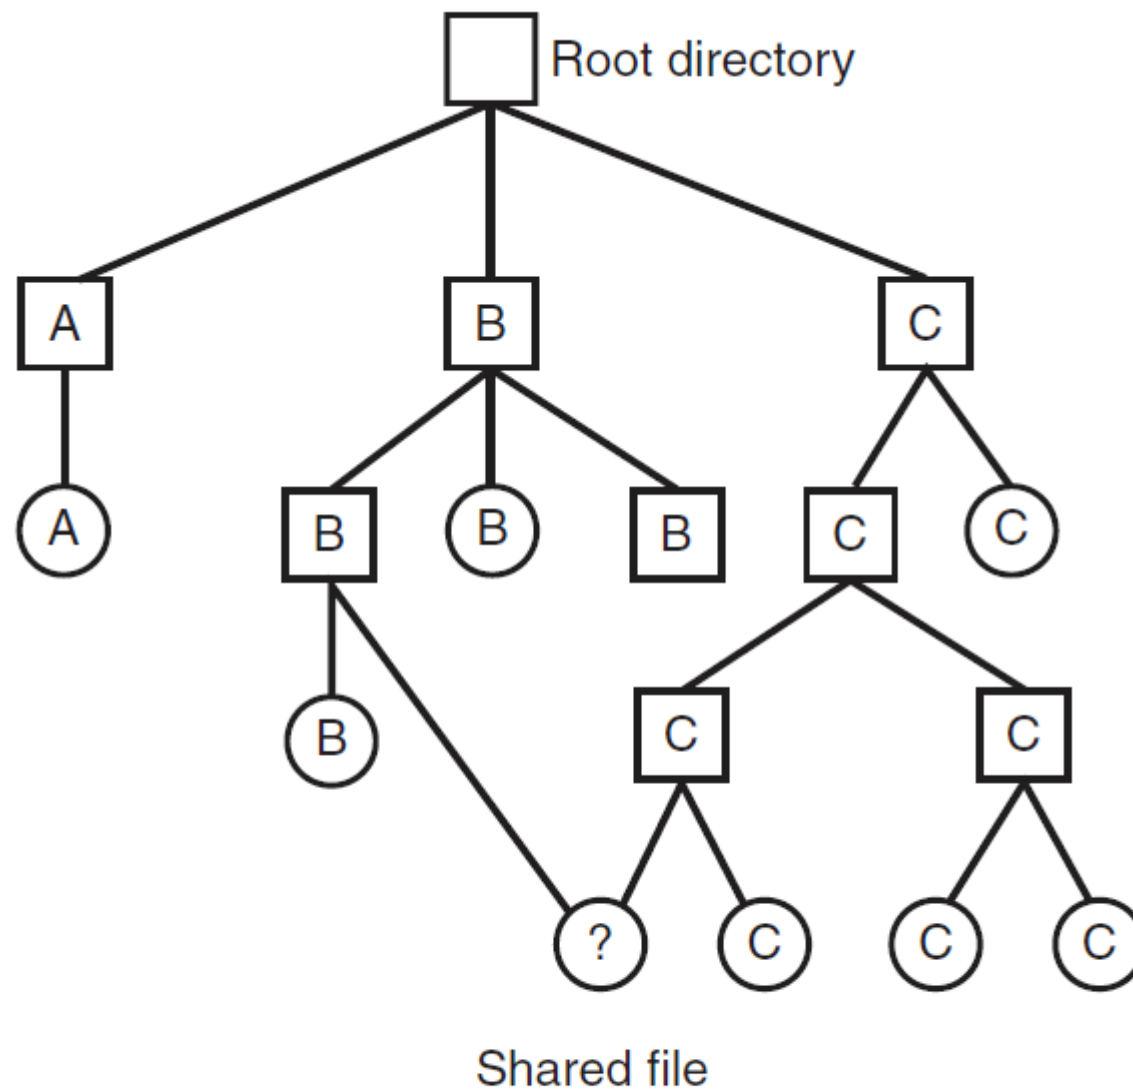


Figure 4-16. File system containing a shared file.

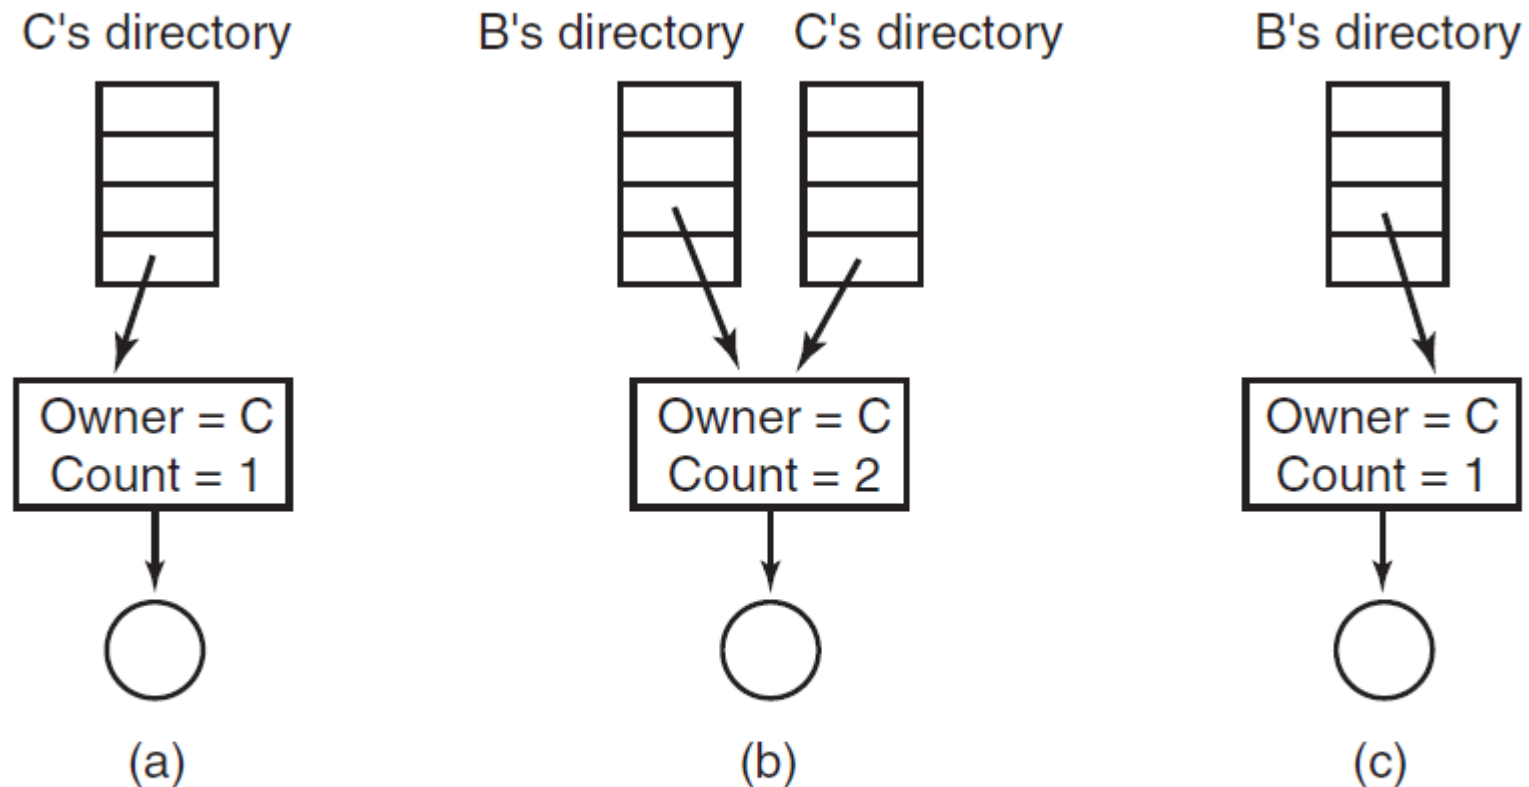


Figure 4-17. (a) Situation prior to linking. (b) After the link is created. (c) After the original owner removes the file.

Directories:

General Graph Directories

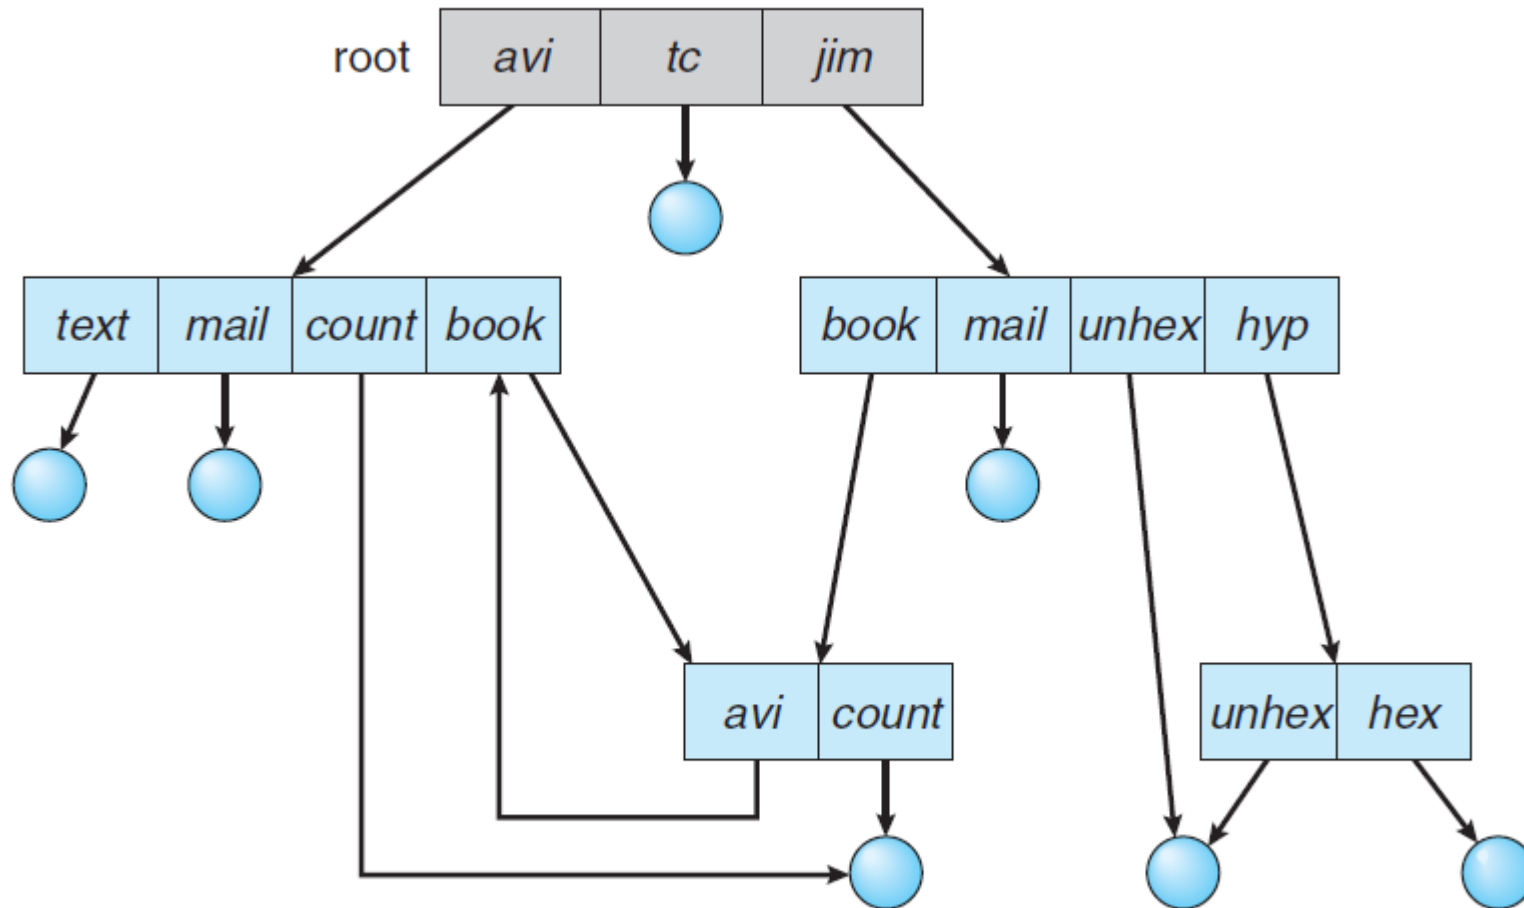


Figure 11.13 General graph directory.

Directory Operations (system calls)

- 1. Create**
- 2. Delete**
- 3. Opendir**
- 4. Closedir**
- 5. Readdir**
- 6. Rename**
- 7. Link**
- 8. Unlink**

File-System Implementation

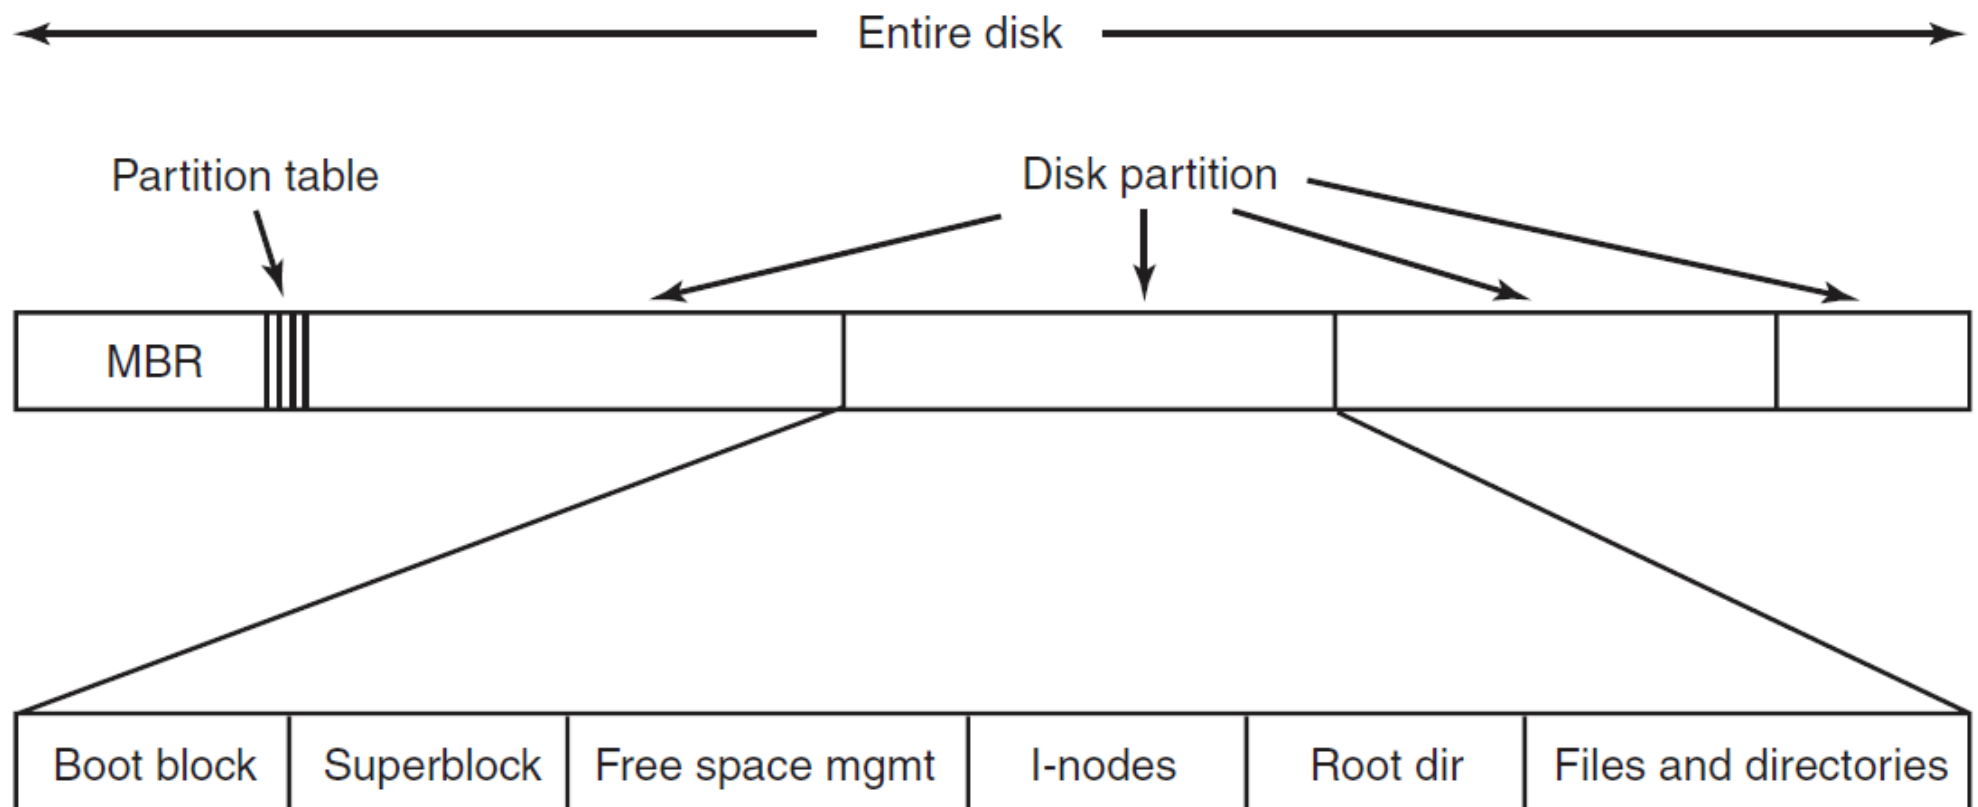


Figure 4-9. A possible file-system layout.

```
# fdisk -l /dev/sda
```

```
Disk /dev/sda: 931,5 GiB, 1000204886016 bytes, 1953525168 sectors
```

```
Units: sectors of 1 * 512 = 512 bytes
```

```
Sector size (logical/physical): 512 bytes / 4096 bytes
```

```
I/O size (minimum/optimal): 4096 bytes / 4096 bytes
```

```
Disklabel type: dos
```

```
Disk identifier: 0xd2cb19de
```

Disposit.	Inicio	Start	Final	Sectores	Size	Id	Tipo
/dev/sda1	*	2048	3074047	3072000	1,5G	7	HPFS/NTFS/exFAT
/dev/sda2		3074048	984686591	981612544	468,1G	7	HPFS/NTFS/exFAT
/dev/sda3		984688638	1927555071	942866434	449,6G	f	W95 Ext'd (LBA)
/dev/sda4		1927555072	1953521663	25966592	12,4G	7	HPFS/NTFS/exFAT
/dev/sda5		984688640	986786296	2097657	1G	83	Linux
/dev/sda6		986787840	1053896703	67108864	32G	82	Linux swap / Solaris
/dev/sda7		1053898752	1927555071	873656320	416,6G	83	Linux

```
# parted /dev/sda
```

```
GNU Parted 3.2
```

```
Usando /dev/sda
```

```
¡Bienvenido/a a GNU Parted! Teclee «help» para ver la lista de órdenes.
```

```
(parted) print
```

```
Modelo: ATA ST1000DM003-1ER1 (scsi)
```

```
Disco /dev/sda: 1000GB
```

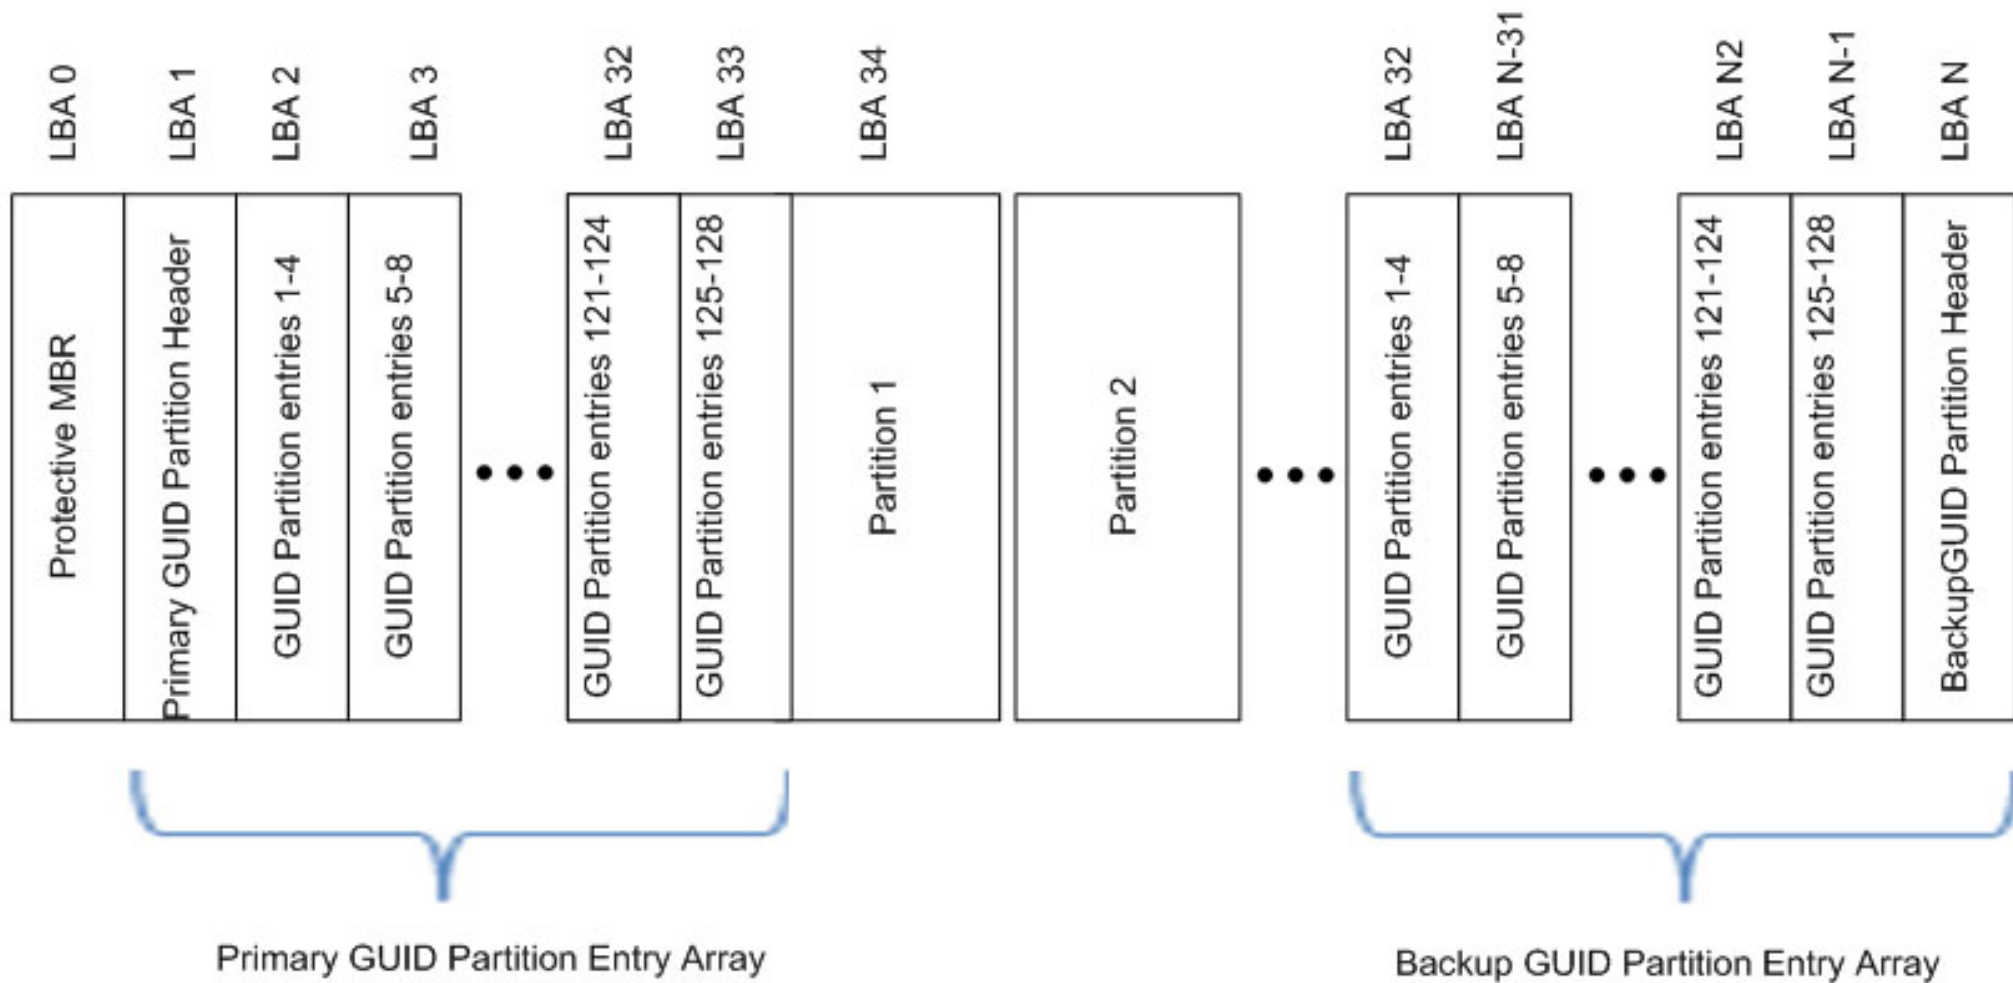
```
Tamaño de sector (lógico/físico): 512B/4096B
```

```
Tabla de particiones: msdos
```

```
Disk Flags:
```

Numero	Inicio	Fin	Tamaño	Tipo	Sistema de archivos	Banderas
1	1049kB	1574MB	1573MB	primary	ntfs	arranque
2	1574MB	504GB	503GB	primary	ntfs	
3	504GB	987GB	483GB	extended		lba
5	504GB	505GB	1074MB	logical	ext4	
6	505GB	540GB	34,4GB	logical	linux-swaps(v1)	
7	540GB	987GB	447GB	logical	ext4	
4	987GB	1000GB	13,3GB	primary	ntfs	

```
(parted) quit
```



<http://macrorit.com/help/partition-gpt-disk.html>

```
# parted /dev/sda
```

```
GNU Parted 3.2
```

```
Usando /dev/sda
```

```
¡Bienvenido/a a GNU Parted! Teclee «help» para ver la lista de órdenes.
```

```
(parted) print
```

```
Modelo: HGST HTS 725050A7E630 (scsi)
```

```
Disco /dev/sde: 500GB
```

```
Tamaño de sector (lógico/físico): 512B/512B
```

```
Tabla de particiones: gpt
```

```
Disk Flags:
```

Numero	Inicio	Fin	Tamaño	Sistema de archivos	Nombre	Banderas
1	1049kB	1050MB	1049MB	ntfs		oculta, diag
2	1050MB	1322MB	273MB	fat32	EFI system partition	arranque, es
3	1322MB	1456MB	134MB		Microsoft reserved partition	msftres
4	1456MB	247GB	245GB	ntfs	Basic data partition	msftdata
6	247GB	264GB	17,2GB	linux-swap(v1)		
7	264GB	486GB	222GB			
5	486GB	500GB	13,9GB	ntfs		oculta, diag

```
(parted) quit
```



```
$ diskutil list
```

```
/dev/disk0 (internal, physical):
```

#:	TYPE	NAME	SIZE	IDENTIFIER
0:	GUID_partition_scheme		*251.0 GB	disk0
1:	EFI	EFI	209.7 MB	disk0s1
2:	Apple_CoreStorage	Macintosh HD	250.1 GB	disk0s2
3:	Apple_Boot	Recovery HD	650.1 MB	disk0s3

```
/dev/disk1 (internal, virtual):
```

#:	TYPE	NAME	SIZE	IDENTIFIER
0:	Apple_HFS	Macintosh HD	+249.8 GB	disk1
		Logical Volume on disk0s2		
		7464223A-BED2-4281-90B2-444BA62069D6		
		Unencrypted		

Implementing Files: Contiguous Allocation




Figure 4-10. (a) Contiguous allocation of disk space for seven files. (b) The state of the disk after files *D* and *F* have been removed.

MOS4E

Implementing Files: Contiguous Allocation





Figure 12.5 Contiguous allocation of disk space.

Implementing Files: Contiguous Allocation




File allocation table

File name	Start block	Length
File A	2	3
File B	9	5
File C	18	8
File D	30	2
File E	26	3

Figure 12.9 Contiguous File Allocation

Implementing Files: Contiguous Allocation



File allocation table

File name	Start block	Length
File A	0	3
File B	3	5
File C	8	8
File D	19	2
File E	16	3

Figure 12.10 Contiguous File Allocation (After Compaction)

Implementing Files: Linked-List Allocation

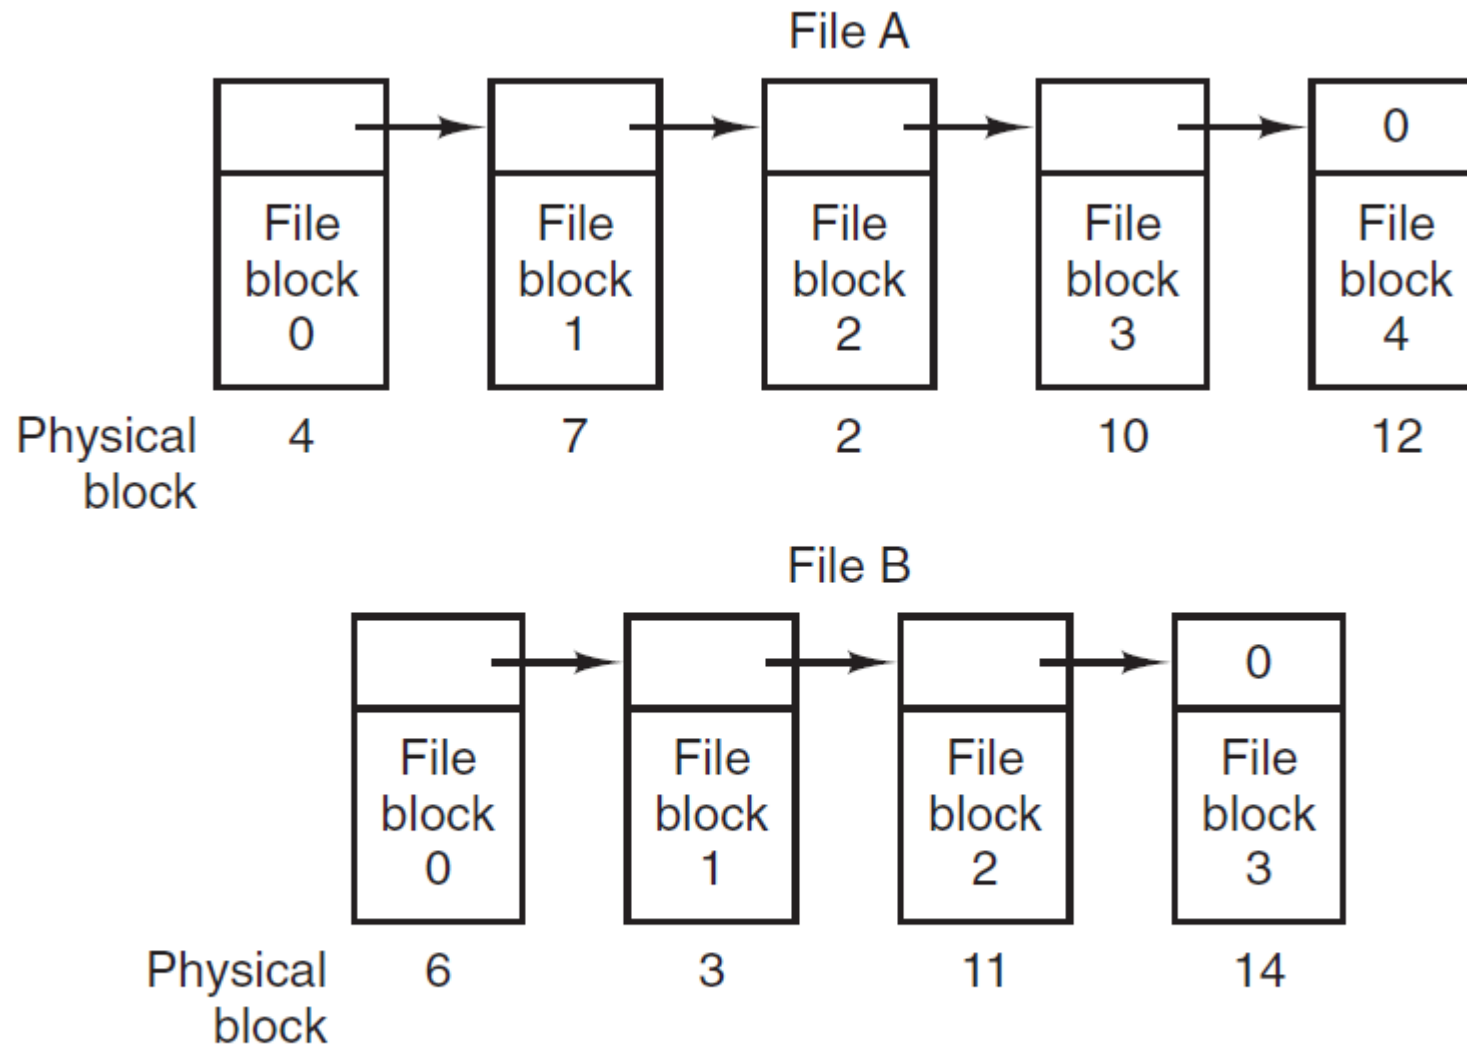


Figure 4-11. Storing a file as a linked list of disk blocks.

Implementing Files: Linked-List Allocation


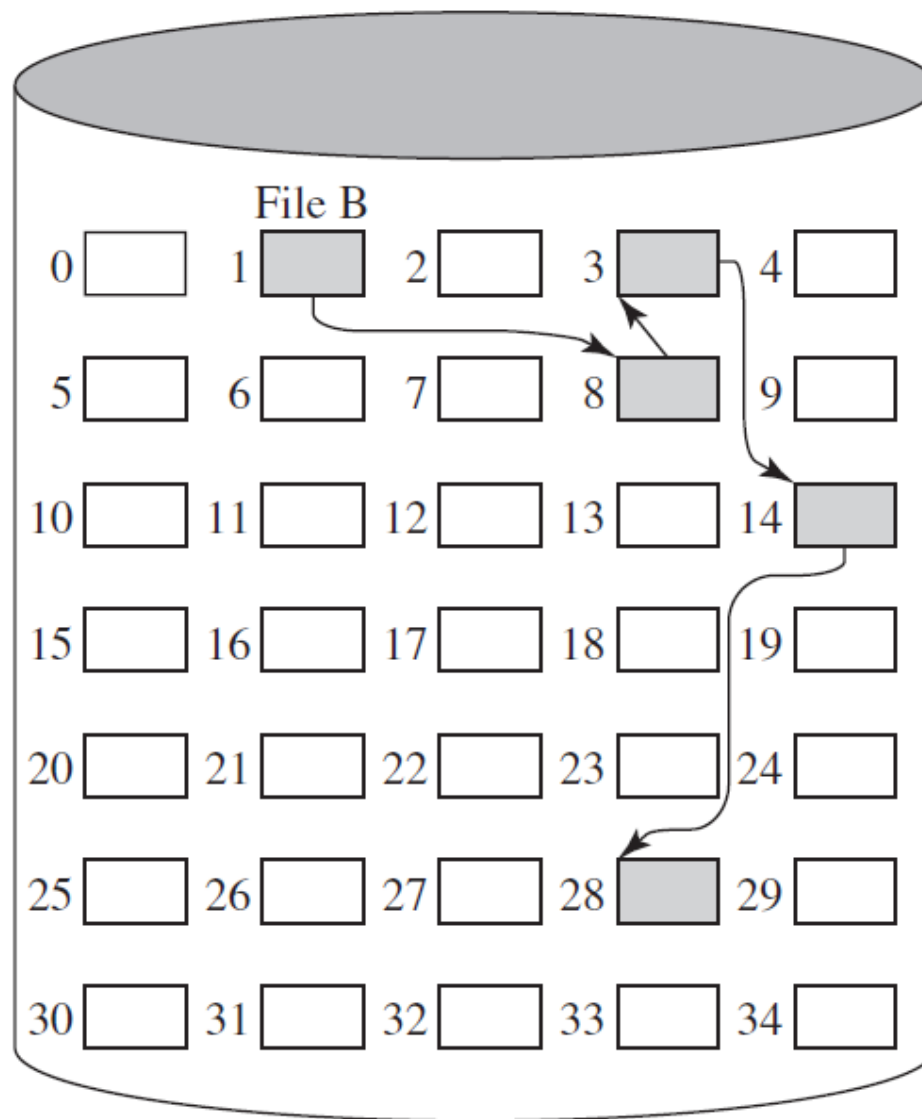


Figure 12.6 Linked allocation of disk space.

Implementing Files: Linked-List Allocation

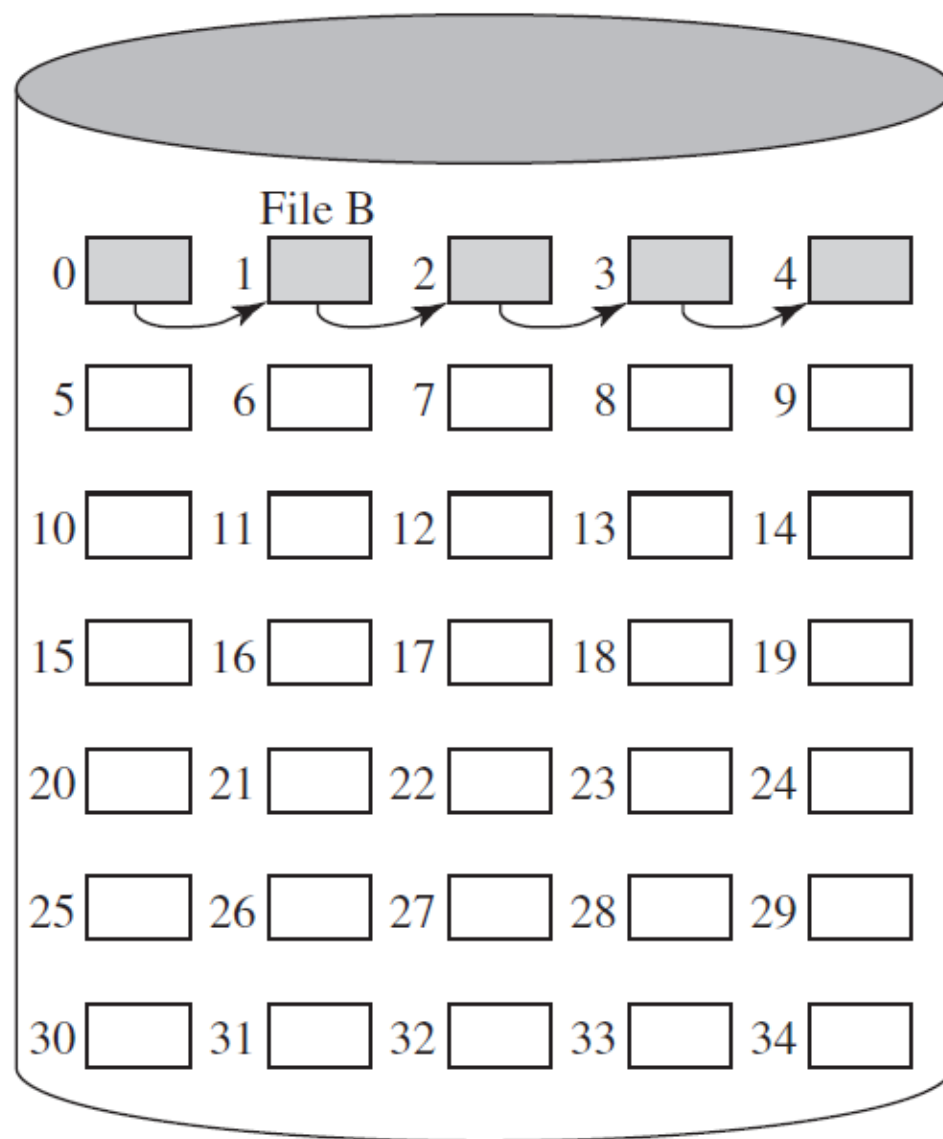


File allocation table

File name	Start block	Length
...
File B	1	5
...

Figure 12.11 Chained Allocation

Implementing Files: Linked-List Allocation



File allocation table

File name	Start block	Length
• • •	• • •	• • •
File B	0	5
• • •	• • •	• • •

Figure 12.12 Chained Allocation (After Consolidation)

Implementing Files:

Linked-List Allocation Using a Table in Memory

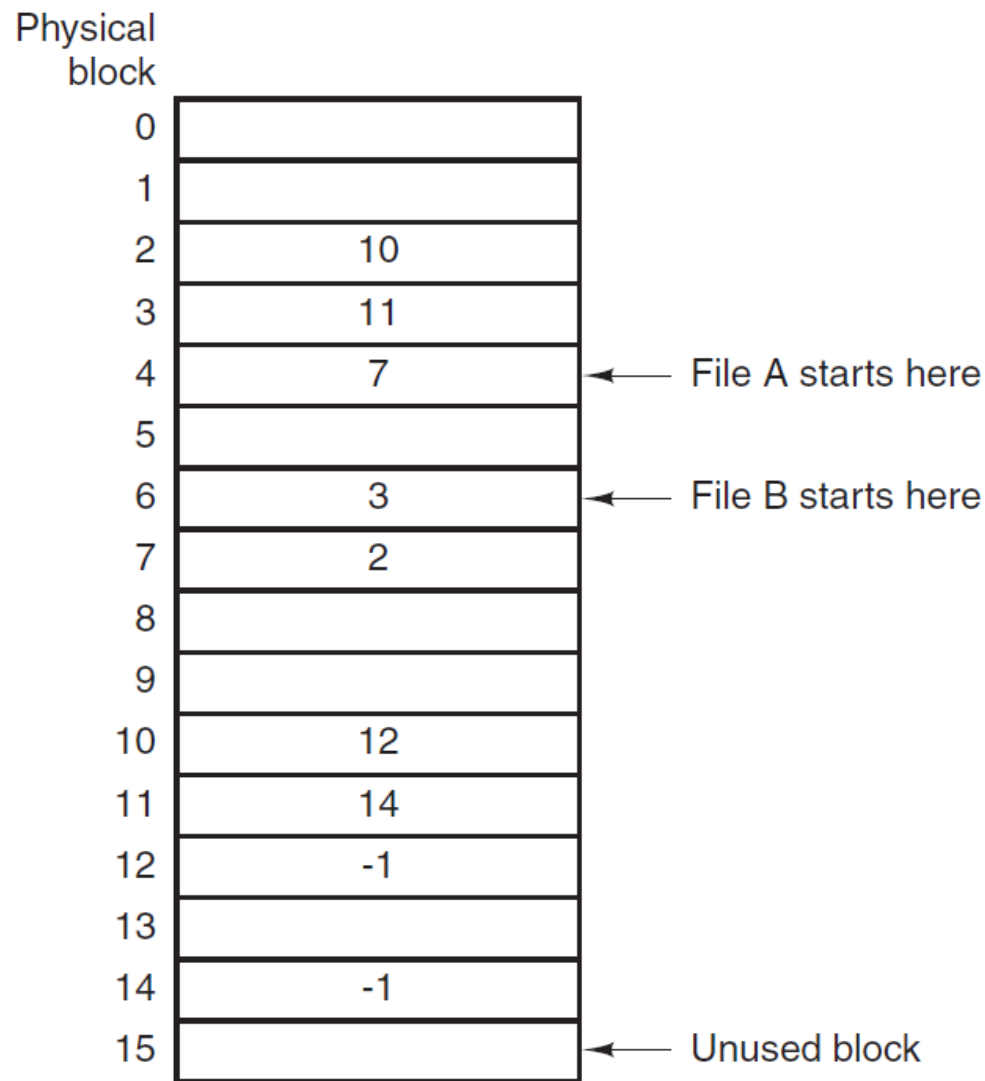


Figure 4-12. Linked-list allocation using a file-allocation table in main memory.

Implementing Files: Linked-List Allocation Using a Table in Memory

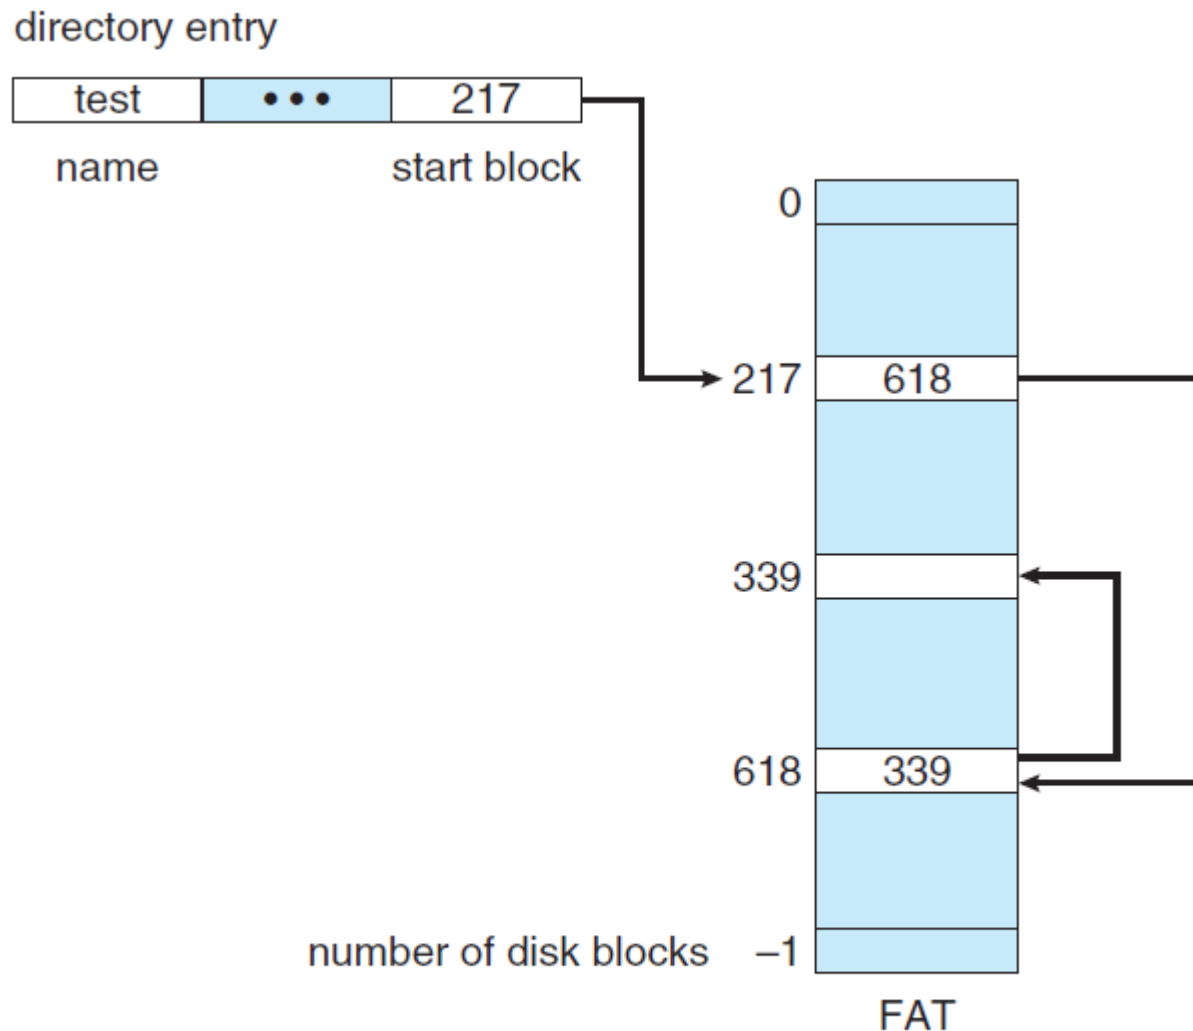


Figure 12.7 File-allocation table.

Implementing Files: Indexed Allocation

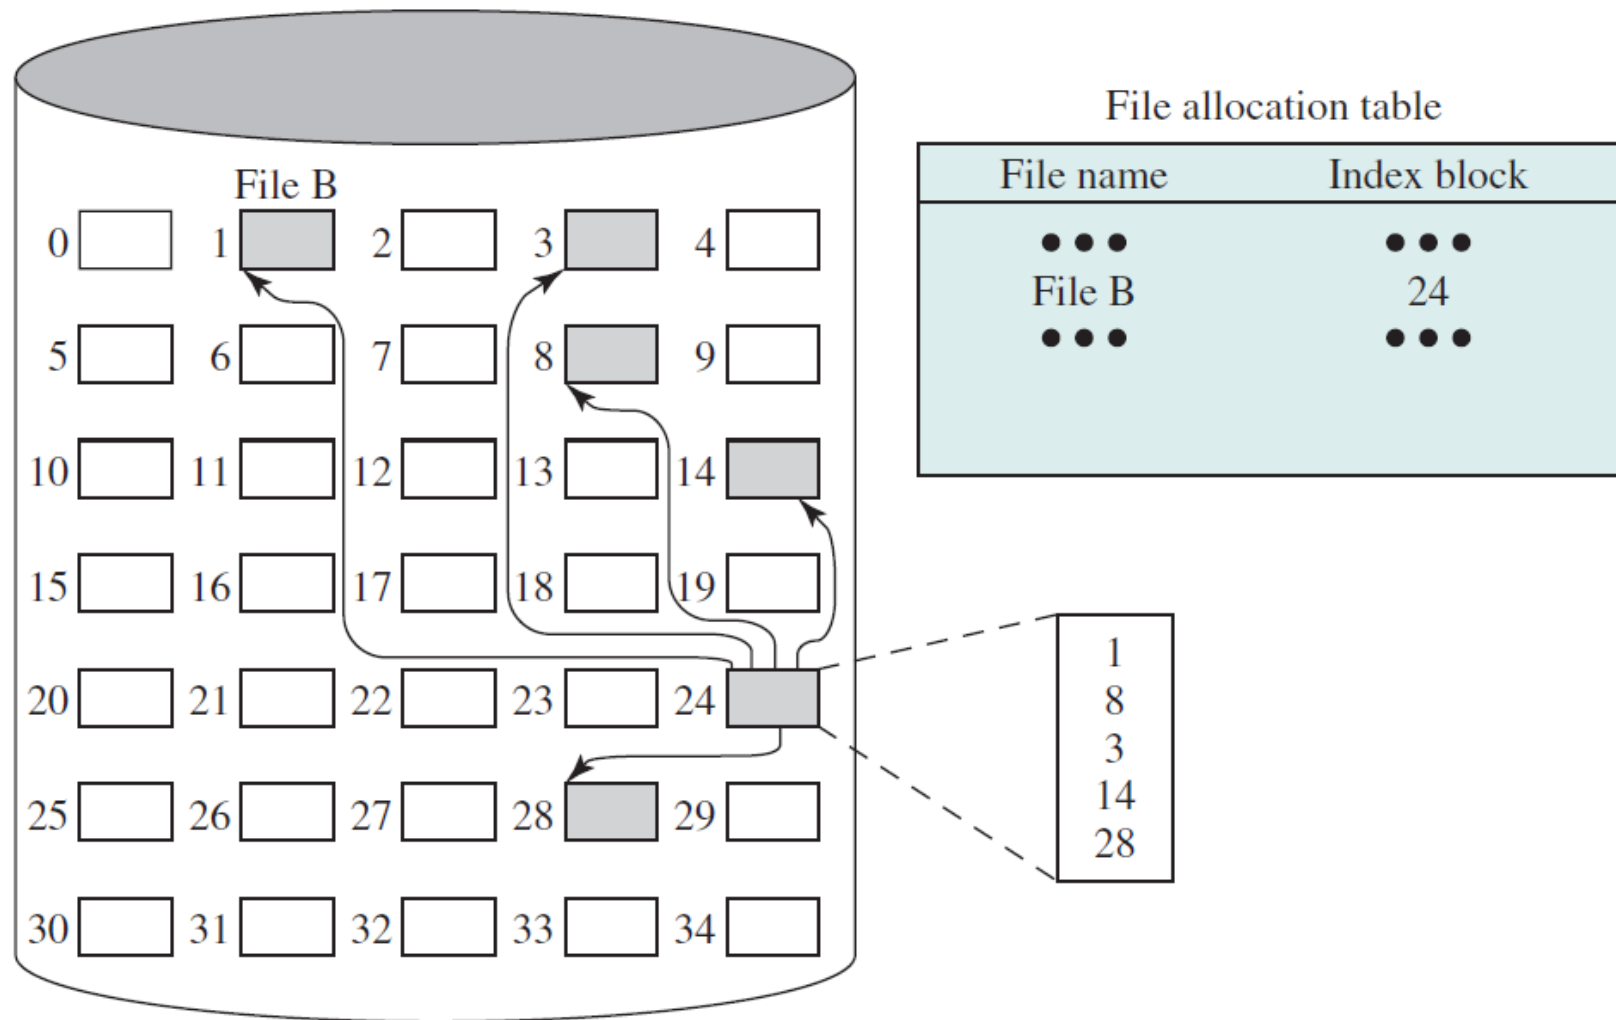


Figure 12.13 Indexed Allocation with Block Portions

Implementing Files: Indexed Allocation

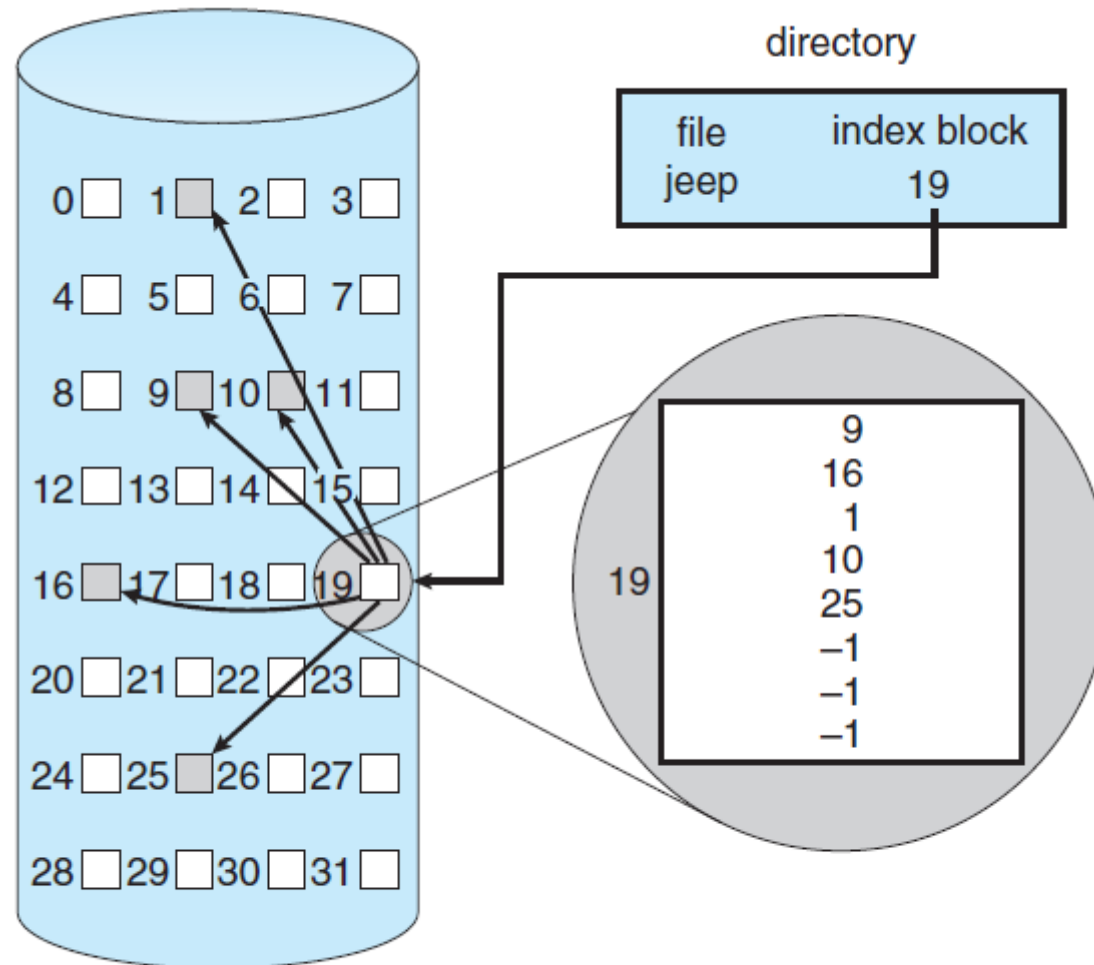


Figure 12.8 Indexed allocation of disk space.

Implementing Files: Indexed Allocation

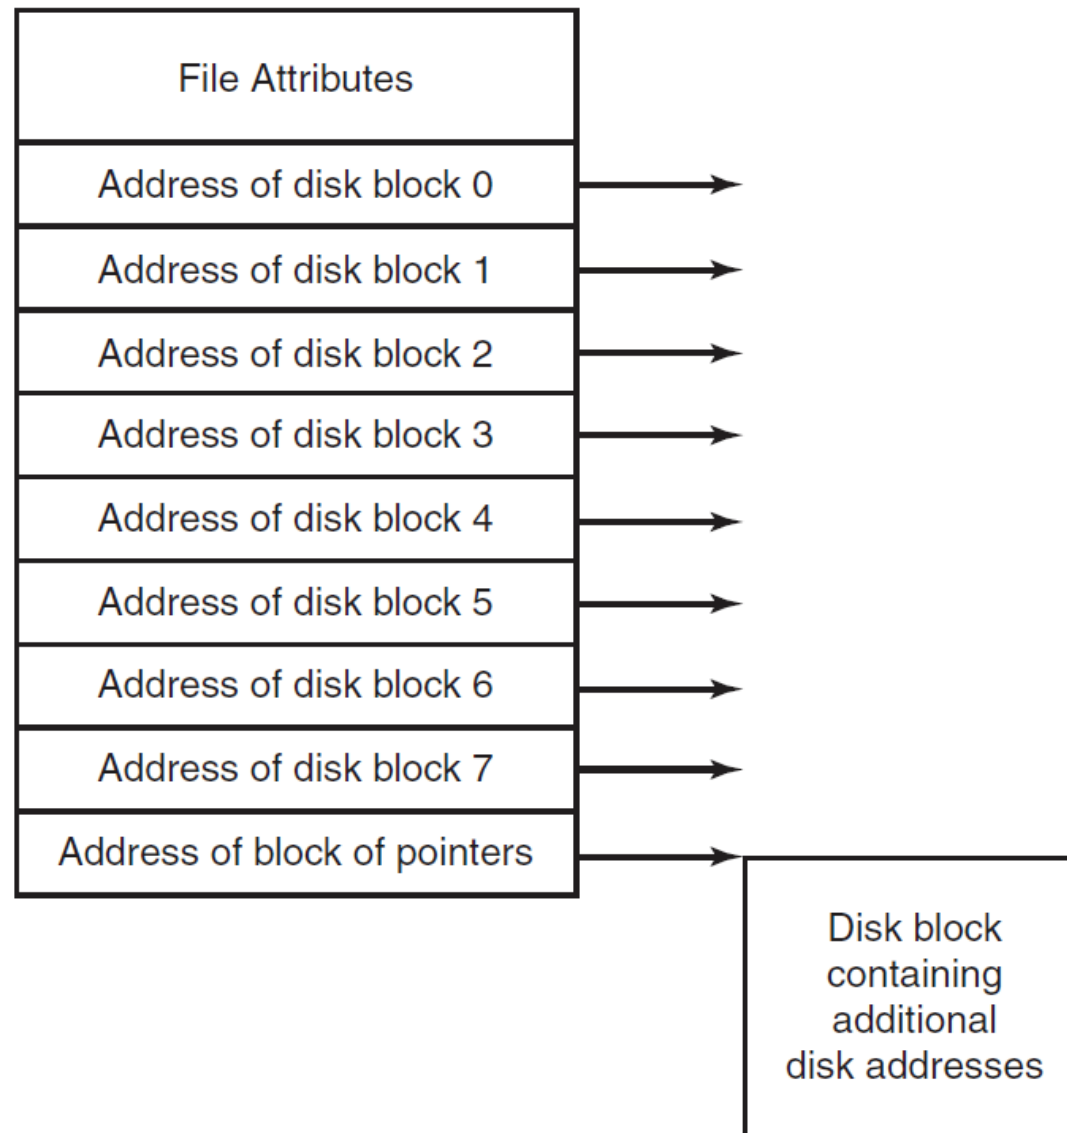


Figure 4-13. An example i-node.

Implementing Files: Indexed Allocation

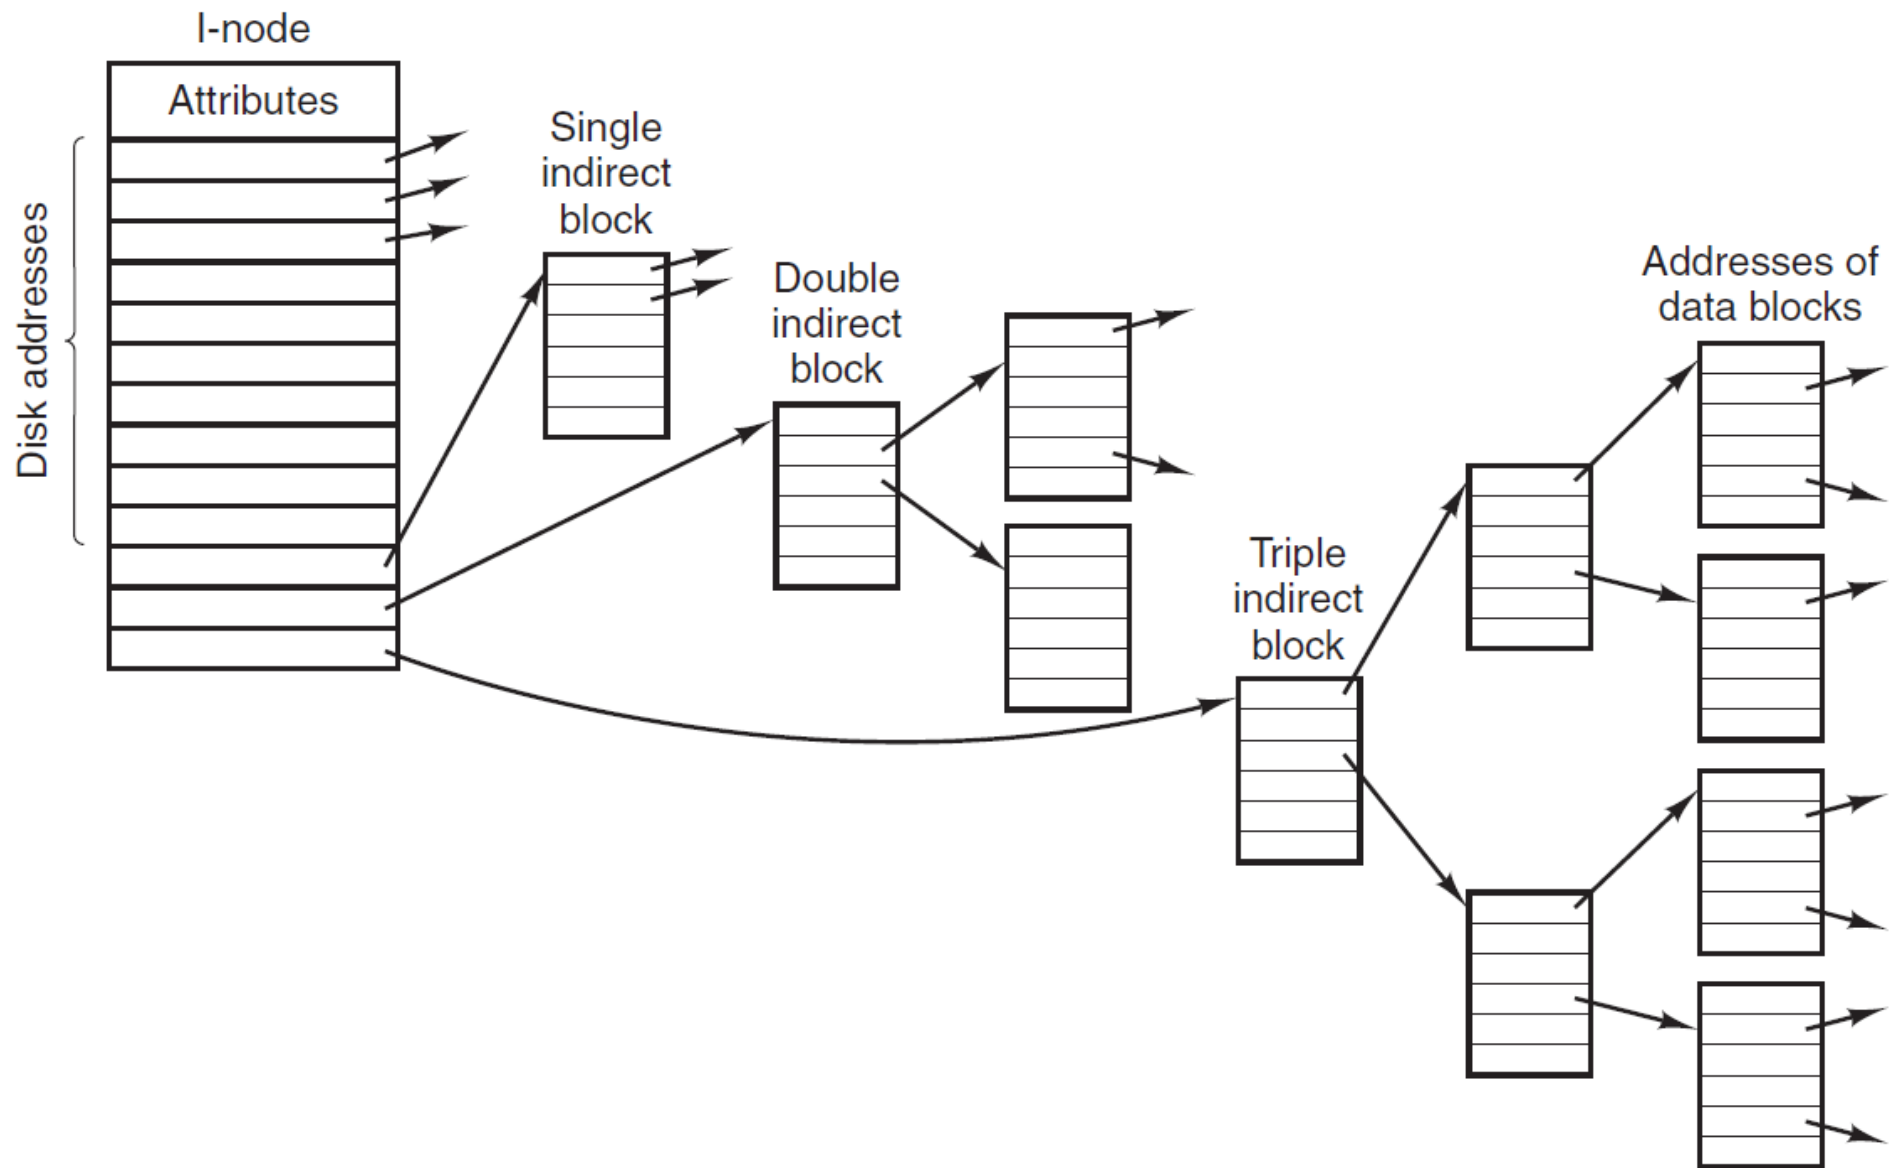


Figure 4-33. A UNIX i-node.

MOS4E

Implementing Files: Indexed Allocation

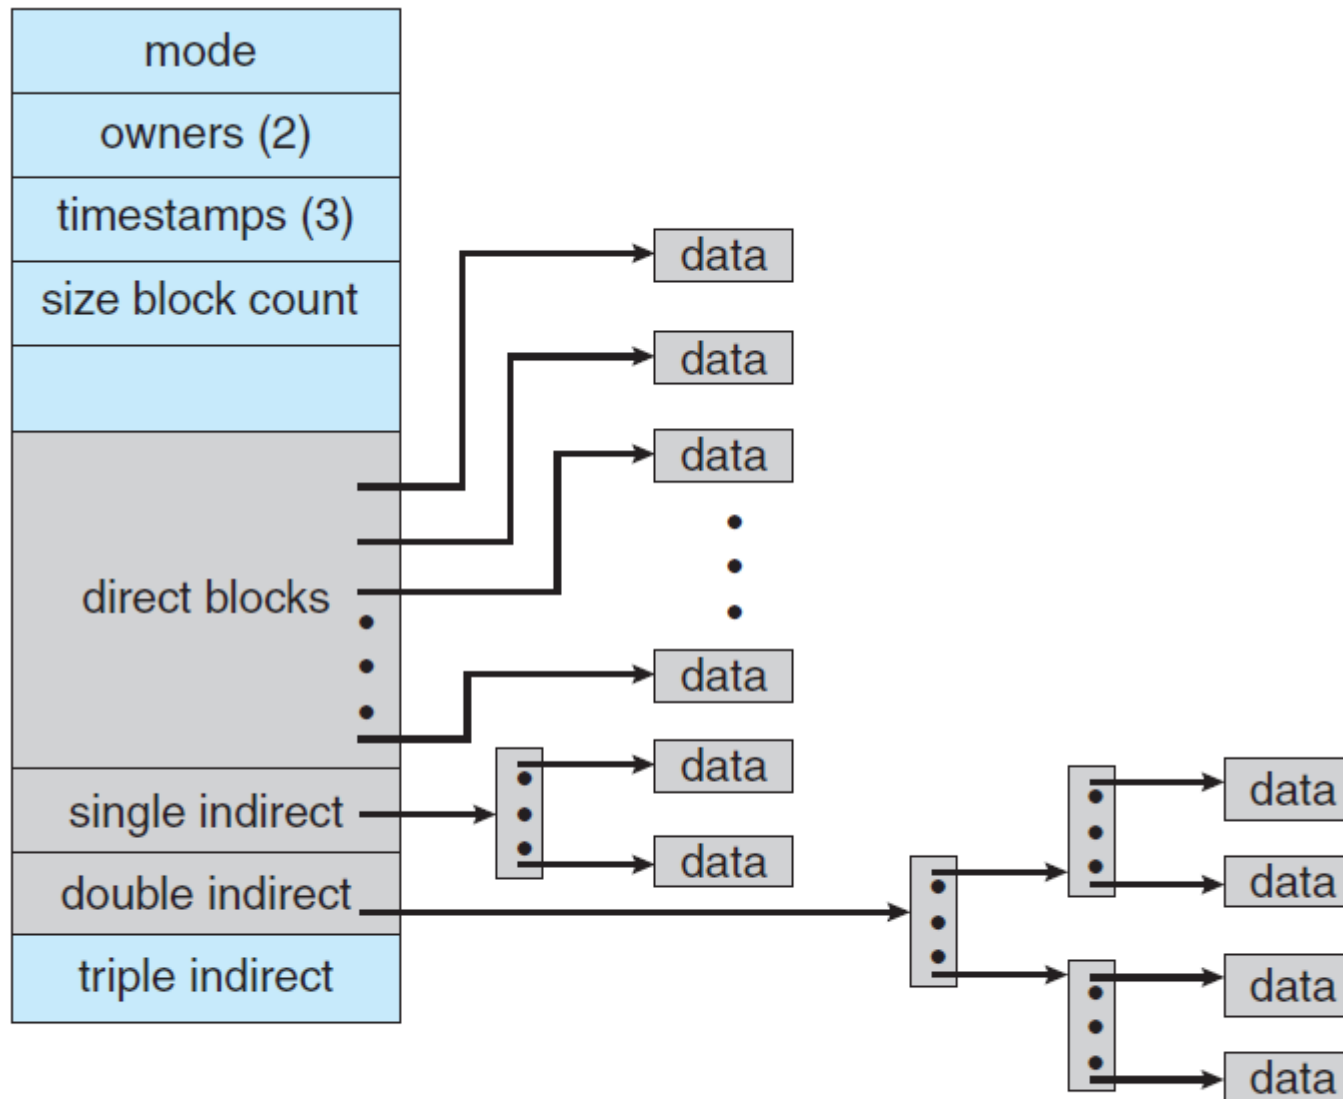



Figure 12.9 The UNIX inode.

OSC9E

“This is the house that Jack built ...”

from Mother Goose




This is the house that Jack built.

This is the malt
That lay in the house that Jack built.

This is the rat,
That ate the malt
That lay in the house that Jack built.


This is the cat,
That killed the rat,
That ate the malt
That lay in the house that Jack built.




Randolf Caldecott

“This is the house that Jack built ...”

from Mother Goose



This is the dog,
That worried the cat,
That killed the rat,
That ate the malt
That lay in the house that Jack built.




This is the cow with the crumpled horn,
That tossed the dog,
That worried the cat,
That killed the rat,
That ate the malt
That lay in the house that Jack built.




Randolf Caldecott

“This is the house that Jack built ...”

from Mother Goose



This is the maiden all forlorn,
That milked the cow with the crumpled horn,
That tossed the dog,
That worried the cat,
That killed the rat,
That ate the malt
That lay in the house that Jack built.




This is the man all tattered and torn,
That kissed the maiden all forlorn,
That milked the cow with the crumpled horn,
That tossed the dog,
That worried the cat,
That killed the rat,
That ate the malt
That lay in the house that Jack built.



Randolf Caldecott

“This is the house that Jack built ...”

from Mother Goose




This is the priest all shaven and shorn,
That married the man all tattered and torn,
That kissed the maiden all forlorn,
That milked the cow with the crumpled horn,
That tossed the dog,
That worried the cat,
That killed the rat,
That ate the malt
That lay in the house that Jack built.




Randolf Caldecott

“This is the house that Jack built ...”

from Mother Goose




This is the cock that crowed in the morn,
That waked the priest all shaven and shorn,
That married the man all tattered and torn,
That kissed the maiden all forlorn,
That milked the cow with the crumpled horn,
That tossed the dog,
That worried the cat,
That killed the rat,
That ate the malt
That lay in the house that Jack built.



Randolf Caldecott

“This is the house that Jack built ...”

from Mother Goose



This is the farmer sowing his corn,
That kept the cock that crowed in the morn,
That waked the priest all shaven and shorn,
That married the man all tattered and torn,
That kissed the maiden all forlorn,
That milked the cow with the crumpled horn,
That tossed the dog,
That worried the cat,
That killed the rat,
That ate the malt
That lay in the house that Jack built.



Randolf Caldecott

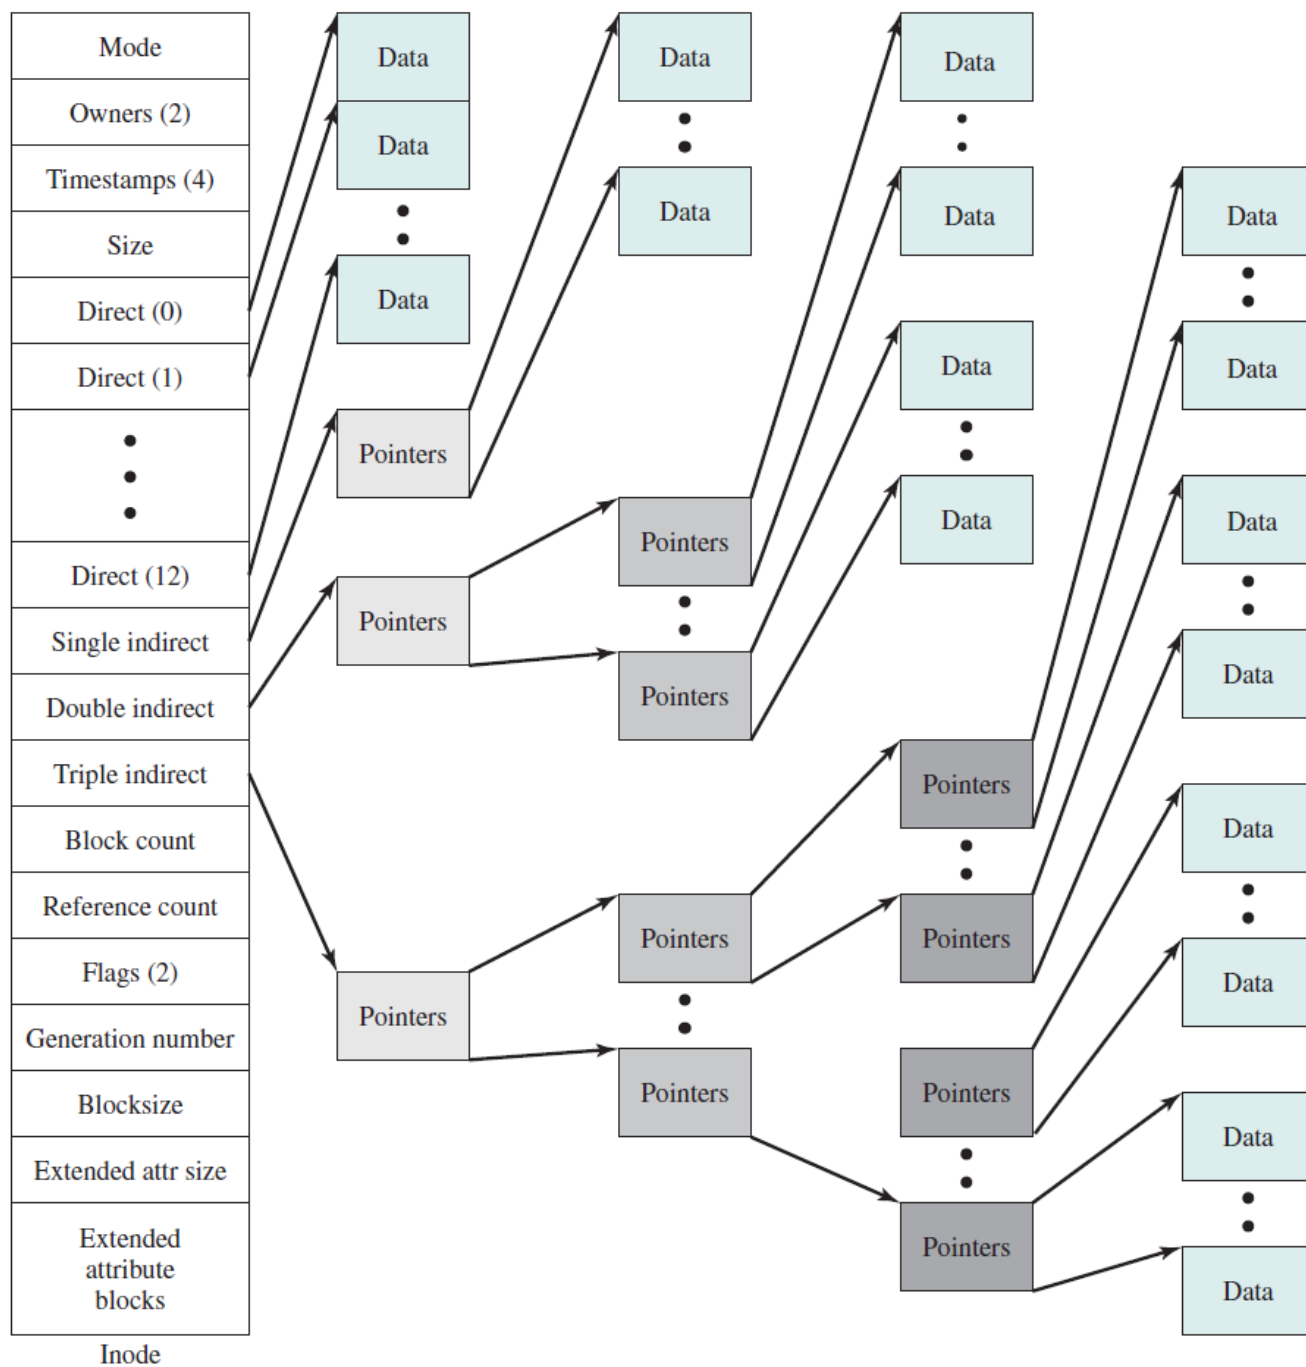


Figure 12.15 Structure of FreeBSD Inode and File

Implementing Directories

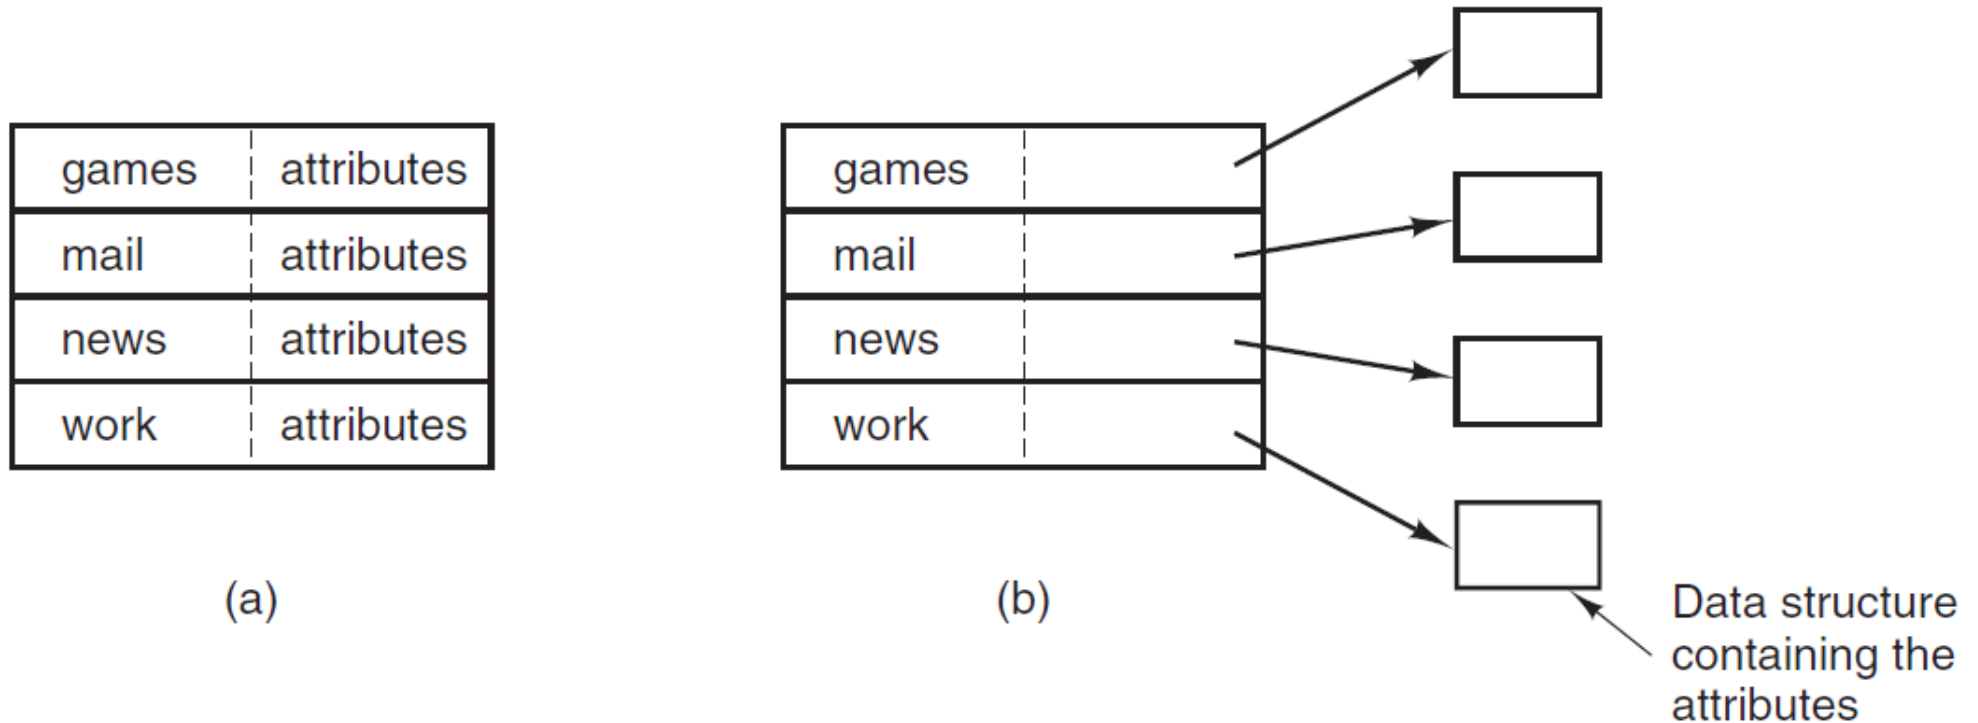


Figure 4-14. (a) A simple directory containing fixed-size entries with the disk addresses and attributes in the directory entry. (b) A directory in which each entry just refers to an i-node.

Implementing Directories




Figure 12.16 UNIX Directories and Inodes

Implementing Directories




Figure 4-32. A UNIX V7 directory entry.

Virtual File Systems

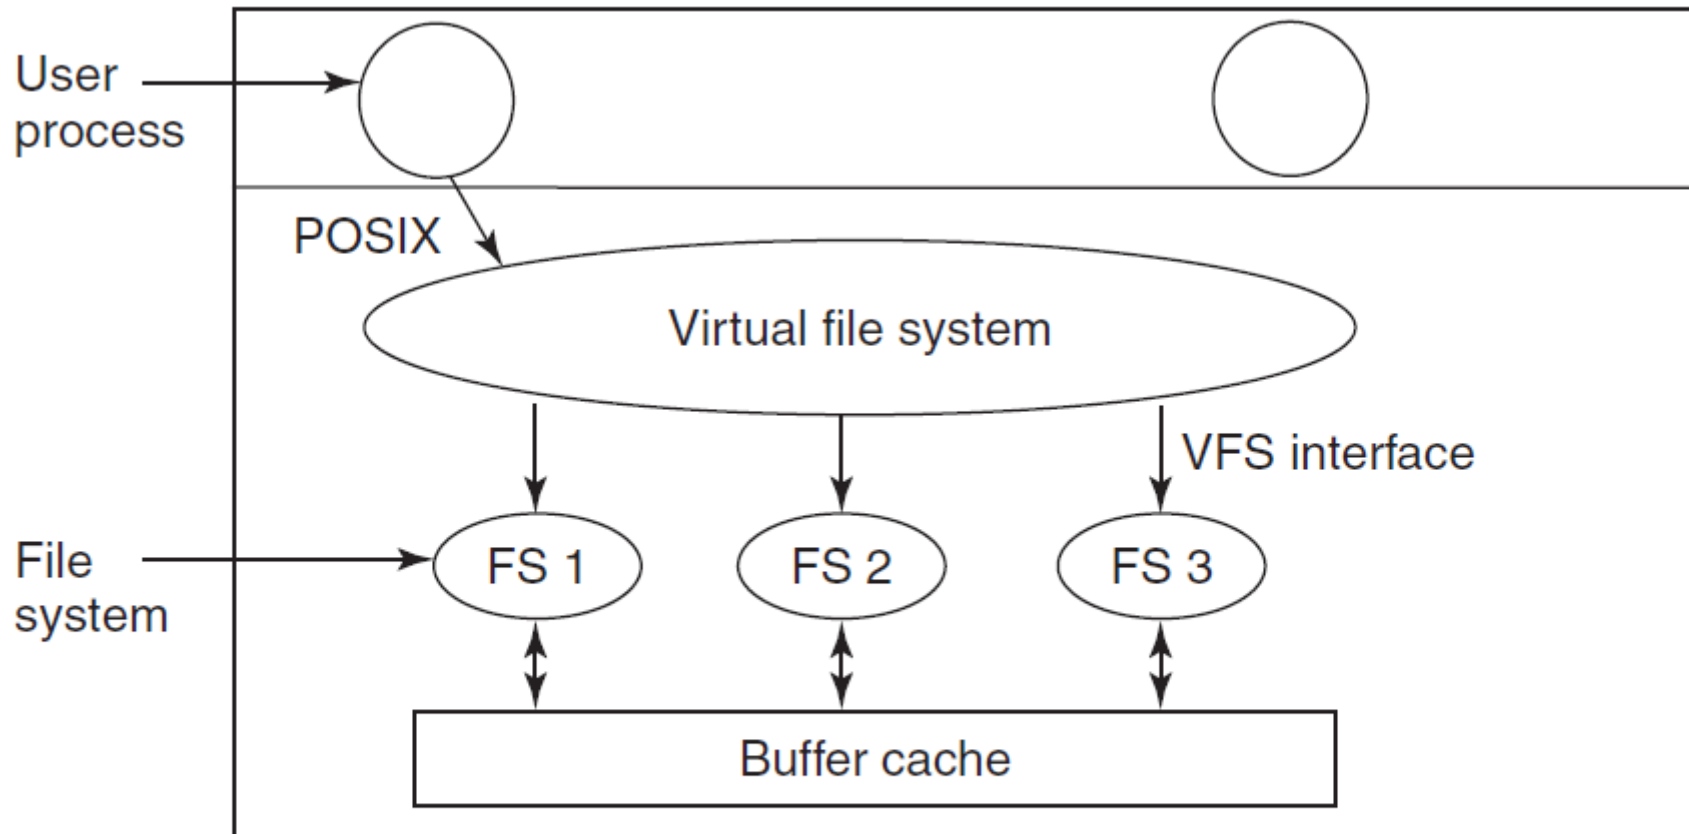


Figure 4-18. Position of the virtual file system.

Virtual File Systems

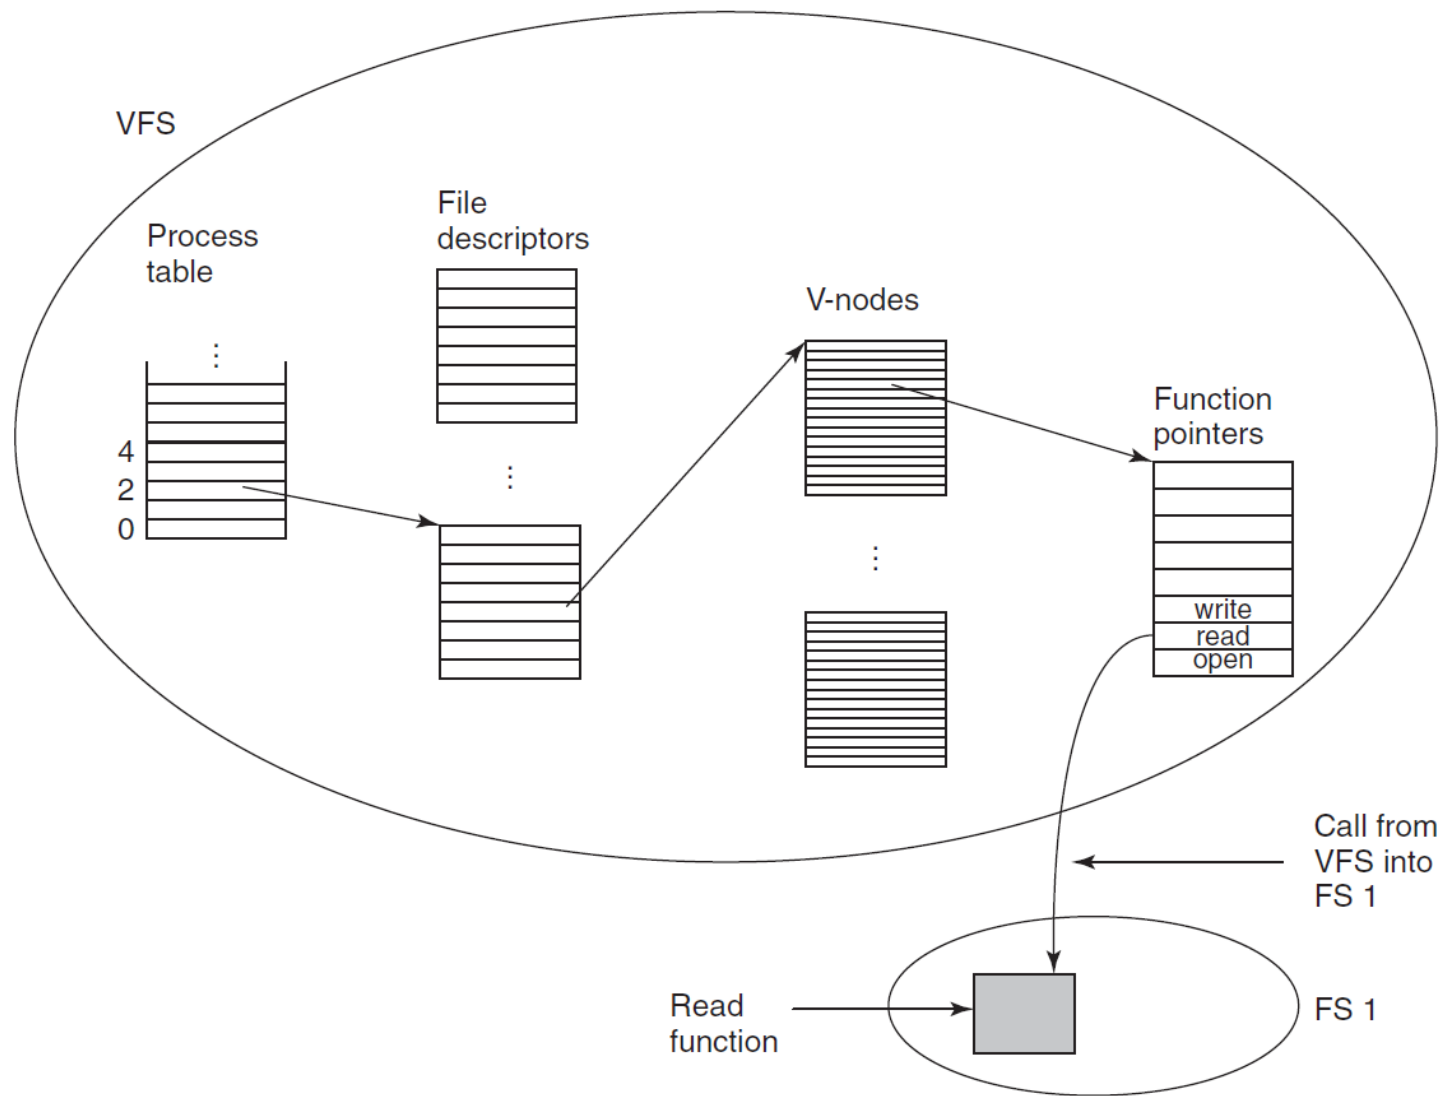


Figure 4-19. A simplified view of the data structures and code used by the VFS and concrete file system to do a read.

Virtual File Systems

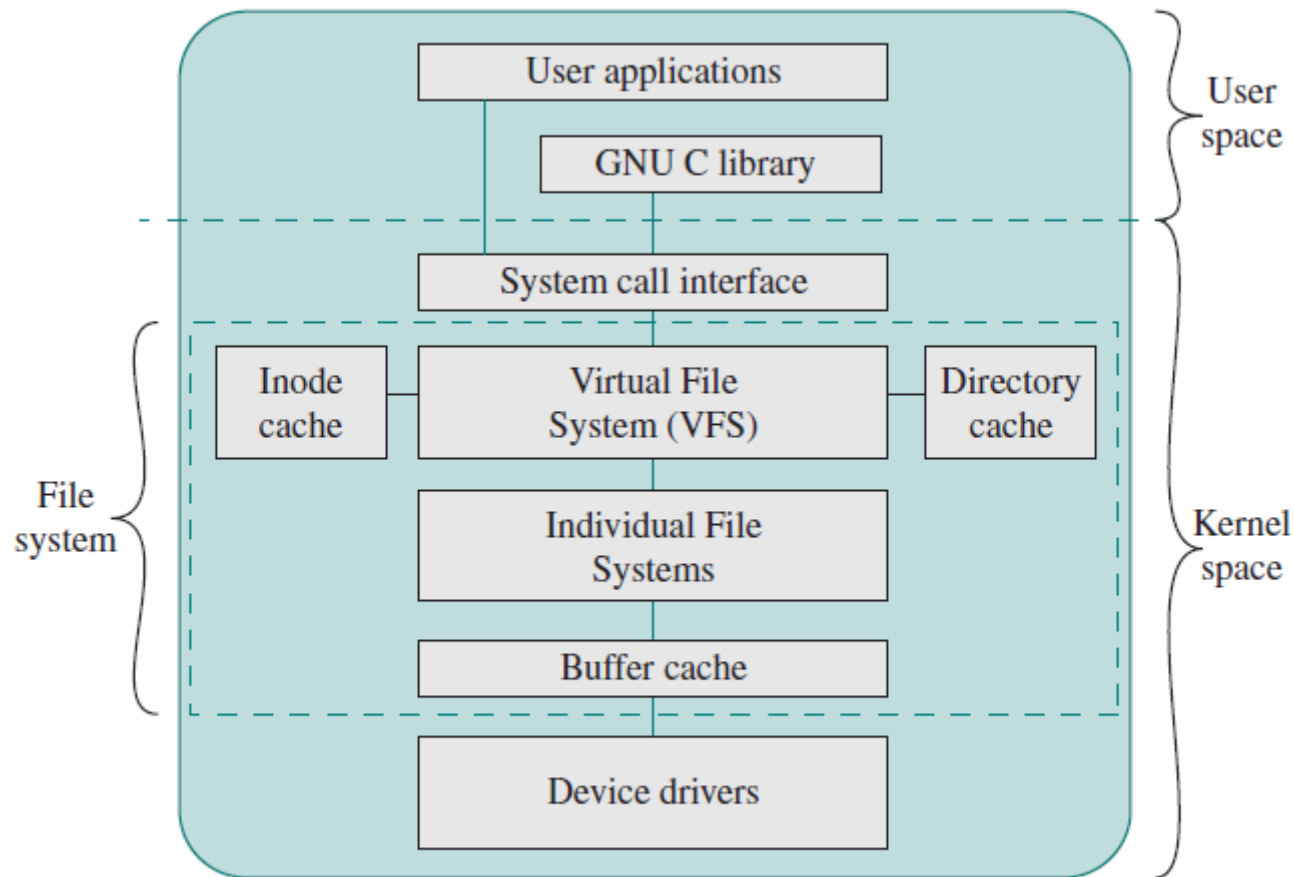


Figure 12.17 Linux Virtual File System Context

Virtual File Systems

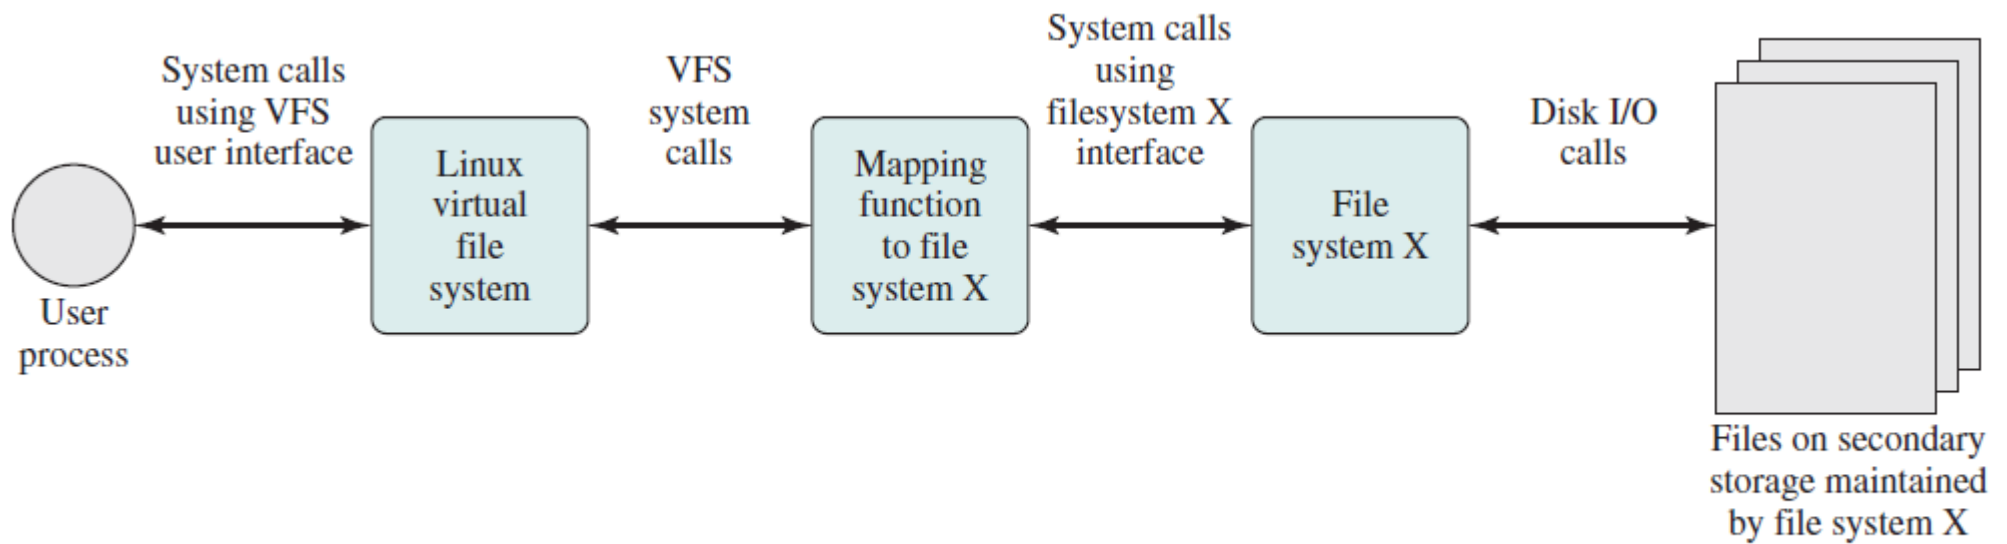


Figure 12.18 Linux Virtual File System Concept

Keeping Track of Free Blocks

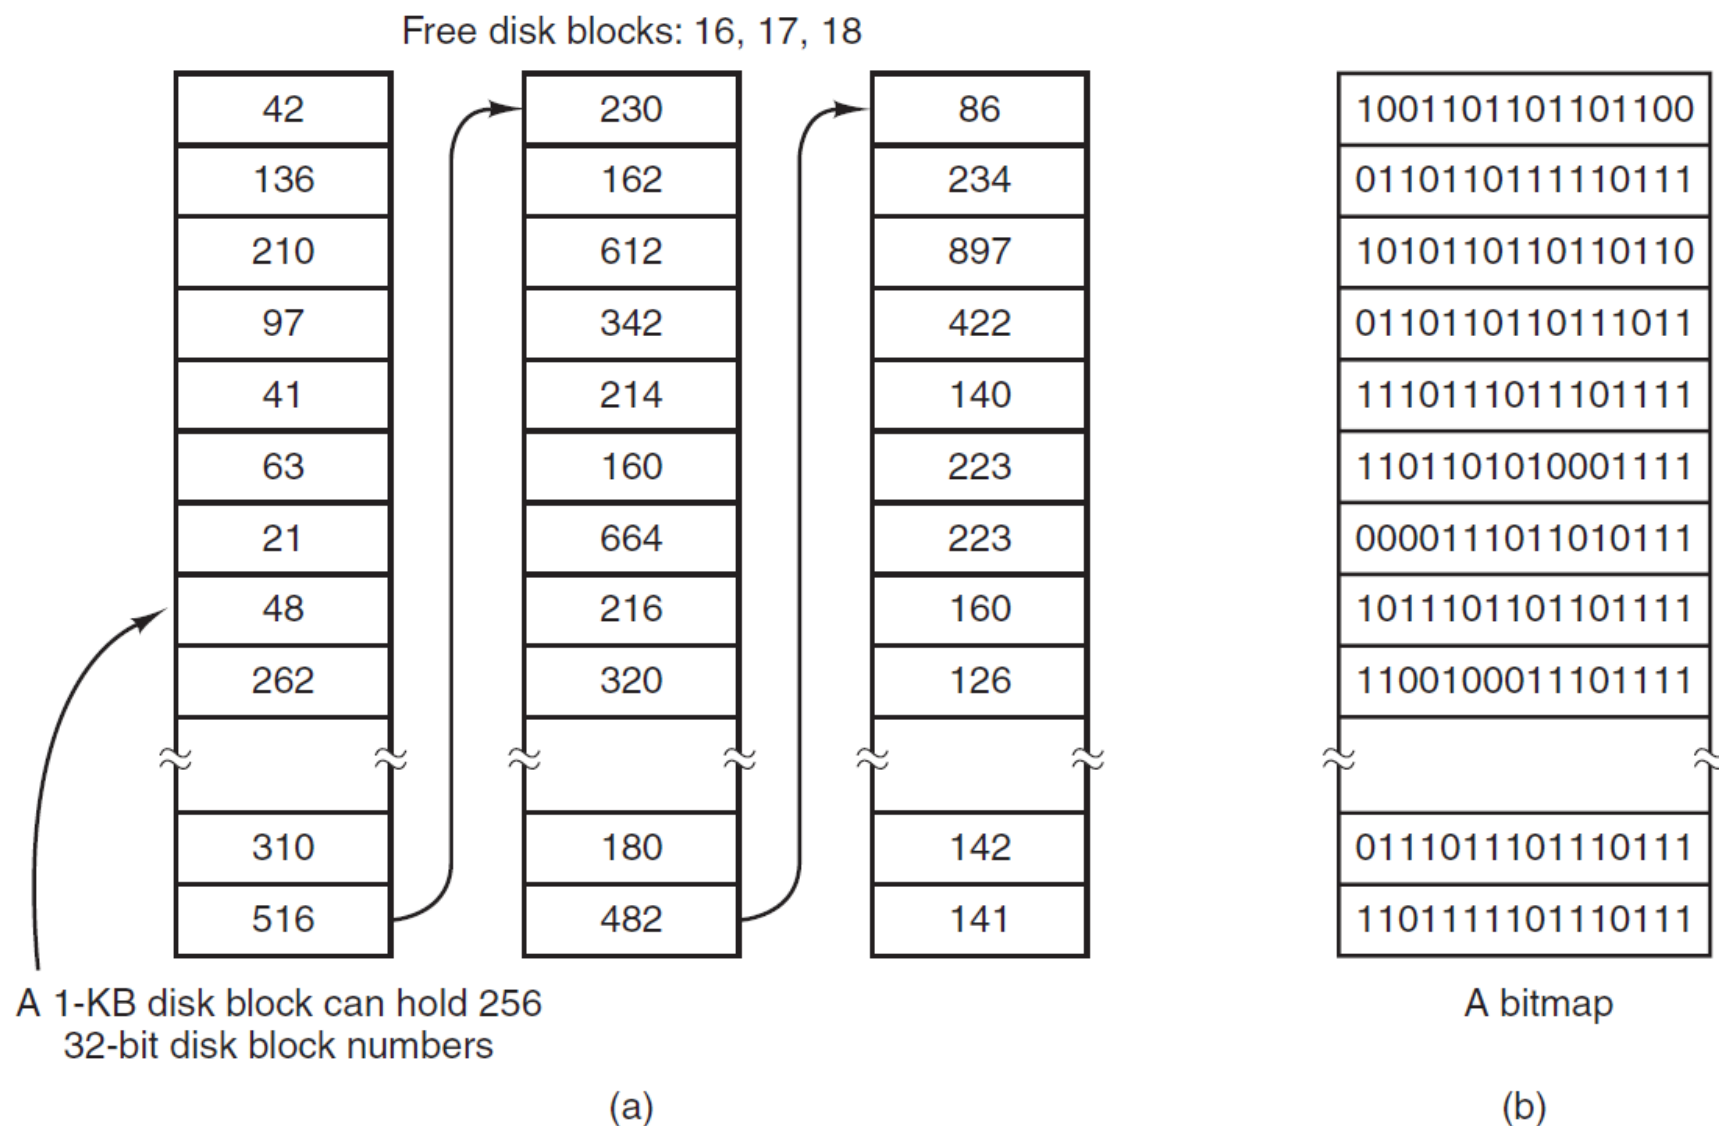


Figure 4-22. (a) Storing the free list on a linked list. (b) A bitmap.

Fast File System

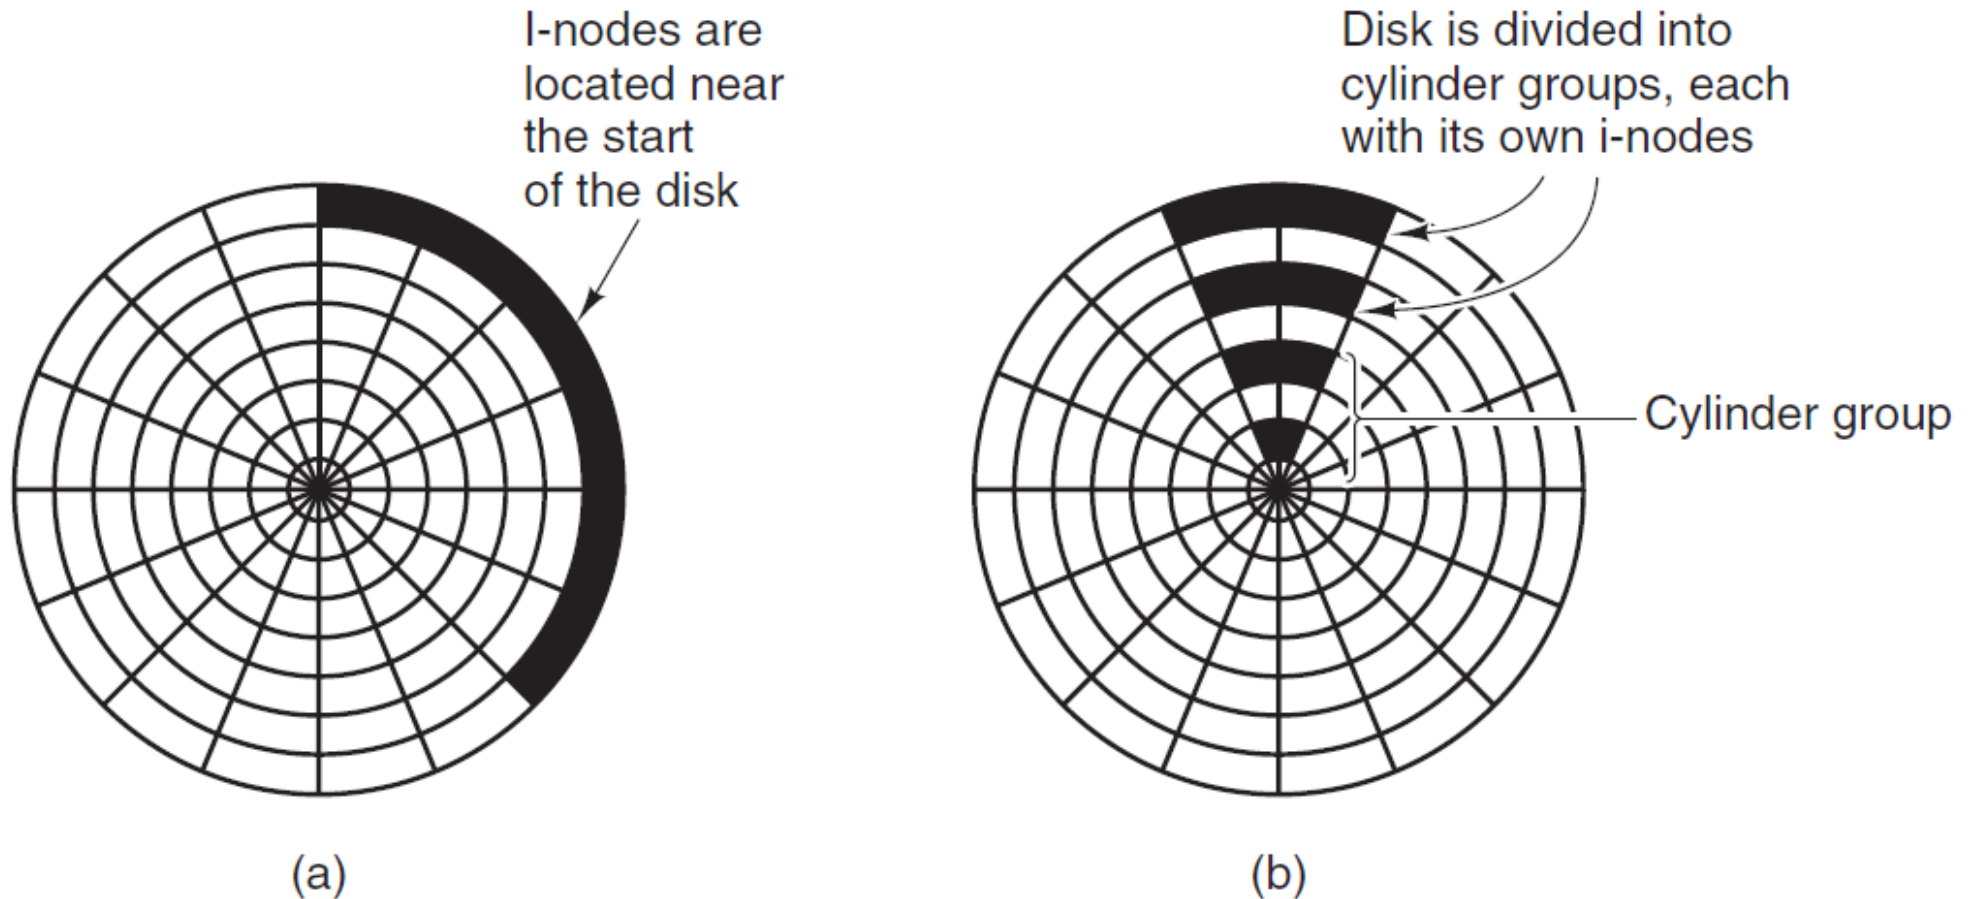


Figure 4-29. (a) I-nodes placed at the start of the disk. (b) Disk divided into cylinder groups, each with its own blocks and i-nodes.

The Linux Ext2 File System

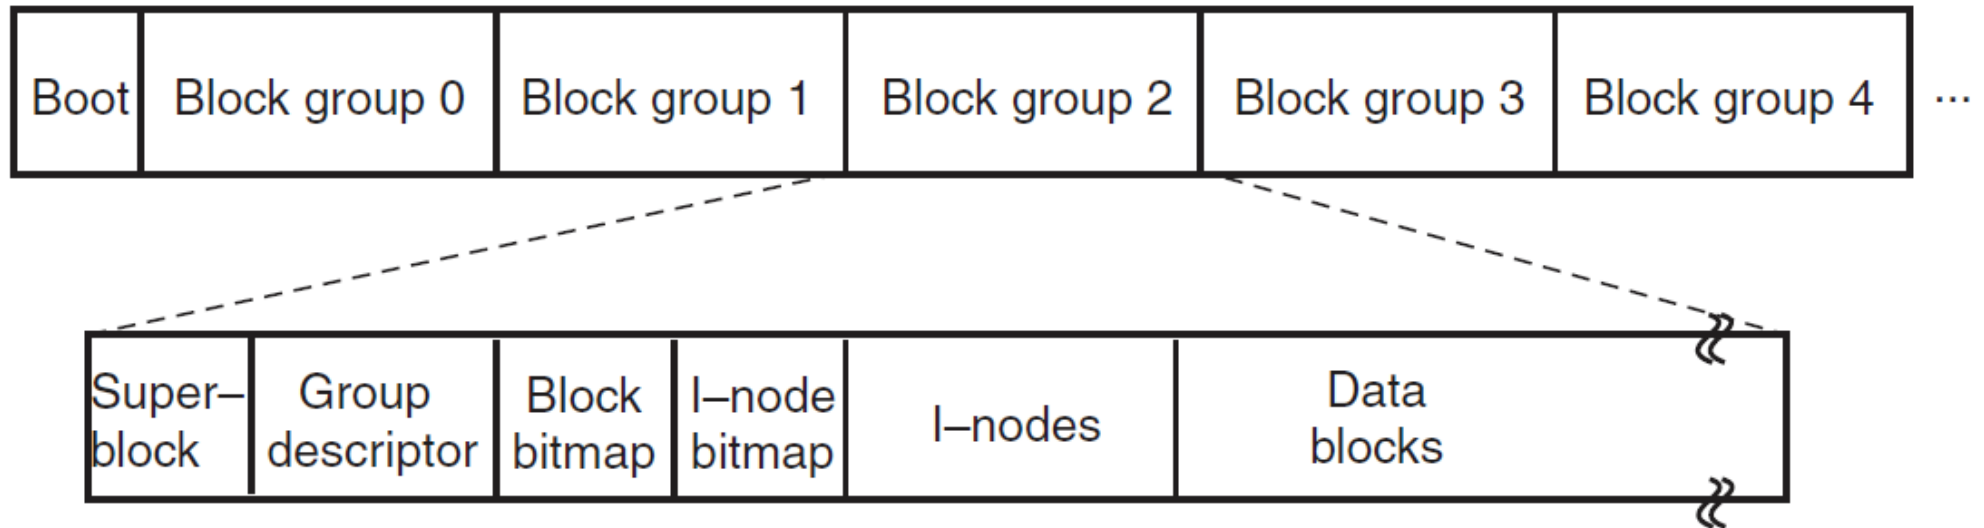


Figure 10-31. Disk layout of the Linux ext2 file system.

The Linux Ext2 File System




Figure 10-32. (a) A Linux directory with three files. (b) The same directory after the file *voluminous* has been removed.

The Linux Ext2 File System

Field	Bytes	Description
Mode	2	File type, protection bits, setuid, setgid bits
Nlinks	2	Number of directory entries pointing to this i-node
Uid	2	UID of the file owner
Gid	2	GID of the file owner
Size	4	File size in bytes
Addr	60	Address of first 12 disk blocks, then 3 indirect blocks
Gen	1	Generation number (incremented every time i-node is reused)
Atime	4	Time the file was last accessed
Mtime	4	Time the file was last modified
Ctime	4	Time the i-node was last changed (except the other times)

Figure 10-33. Some fields in the i-node structure in Linux.